

# Microeconometrics (Causal Inference)

## Week 2 - Linear regression (OLS) and bootstrapping

Joshua D. Merfeld  
KDI School of Public Policy and Management

2023-08-18


## What are we doing today?

- ▶ This week we will review linear regression (OLS) and its assumptions
  - ▶ We will also discuss how to interpret the results of a regression
- ▶ Some of this will be review, but there will be some new stuff, too
  - ▶ For example, we will discuss bootstrapping on Thursday


## What are we doing today?

- ▶ I will be going through some econometric theory this week
- ▶ But we will also be using some data throughout the week
  - ▶ This data comes from the [World Development Indicators](#) database
  - ▶ I have uploaded the csv file to this week's GitHub folder

## Fitting lines to data



## Which line best fits the data?



## Ordinary Least Squares (OLS)

- ▶ OLS is a method for fitting a line to data
  - ▶ It is the most common method for fitting lines to data
- ▶ You have an outcome, let's call it  $y_i$  (where  $y$  is the outcome and  $i$  is an individual)
- ▶ You also have a set of covariates, let's call them  $x_{ik}$  (where  $k$  denotes different covariates)
- ▶ The relationship might look something like this, assuming linearity:

$$y_i = \beta_0 + \beta_1 x_{i,1} + \beta_2 x_{i,2} + \cdots + \beta_k x_{i,k} + \epsilon_i \quad (1)$$

- ▶ OLS is about finding the **average** relationship between  $Y$  and a set of covariates,  $X$ 
  - ▶ The fit will never be perfect, which is why we have  $\epsilon_i$
  - ▶  $\epsilon_i$  is the error term and shows that individuals will deviate from the average relationship
  - ▶ For example, Bill Gates is a high-school drop out but is still a billionaire; this is deviation from an average relationship between education and income

## Minimizing the sum of squared errors

- ▶ In order to find a “line of best fit”, we need an objective function
  - ▶ We want to minimize the distance between the line and the data, but how do we define that distance?
- ▶ We can define the distance as the sum of squared errors (SSE)
  - ▶ The SSE is the sum of the squared difference between the actual value and the predicted value
  - ▶ The predicted value is the value on the line of best fit
- ▶ In other words:

$$\text{SSE} = \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (2)$$



$$\text{SSE} = \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (3)$$

- ▶  $\hat{y}_i$  is the predicted value of  $y_i$  for individual  $i$
- ▶ But how do we find  $\hat{y}_i$ ?
  - ▶ We need to find the line of best fit
  - ▶ We need to find the values of  $\beta_0, \beta_1, \dots, \beta_k$  that minimize the SSE.
- ▶ We can do this by noting that  $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_{i,1} + \hat{\beta}_2 x_{i,2} + \dots + \hat{\beta}_k x_{i,k}$ :

$$\text{SSE} = \sum_{i=1}^n (y_i - \beta_0 - \hat{\beta}_1 x_{i,1} - \hat{\beta}_2 x_{i,2} - \dots - \hat{\beta}_k x_{i,k})^2 \quad (4)$$

- ▶ Let's rewrite the equation in matrix form:
- ▶  $\mathbf{Y}$  is a vector of outcomes,  $\mathbf{X}$  is a matrix of covariates, and  $\beta$  is a vector of coefficients:

$$\text{SSE} = (\mathbf{Y} - \mathbf{X}\beta)'(\mathbf{Y} - \mathbf{X}\beta) \quad (5)$$

- ▶ It turns out that this becomes an optimization problem
- ▶ We know the  $\mathbf{Y}$  and  $\mathbf{X}$ , so we want to minimize this with respect to  $\beta$ :

$$\min_{\beta} (\mathbf{Y} - \mathbf{X}\beta)'(\mathbf{Y} - \mathbf{X}\beta) \quad (6)$$

- ▶ We know the solution to this:

$$\hat{\beta} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Y} \quad (7)$$

- ▶ Seer this formula into your brain!

## A problem: uncertainty

- ▶ We have a solution, but we have a problem: uncertainty
  - ▶ We generally do not have the population, which means  $\hat{\beta}$  is just a sample estimate of the population parameters  $\beta$
  - ▶ Note the use of hats to denote sample estimates
- ▶ We need to know how uncertain we are about our estimates
  - ▶ We need to know how much  $\hat{\beta}$  will vary from sample to sample
  - ▶ Before we do this, we need to do a little bit of background on probability and statistics

## Two important theorems

- ▶ There are two important theorems that we need to know about
  - ▶ The Law of Large Numbers
  - ▶ The Central Limit Theorem
  
- ▶ Let's go through these one at a time

- ▶ The LLN says that as the sample size increases, the sample mean will converge to the population mean
  - ▶ In other words, as  $n \rightarrow \infty$ ,  $\bar{x} \rightarrow \mu$ 
    - ▶ where  $n$  is the sample size,  $\bar{x}$  is the sample mean,  $\mu$  is the true population mean
- ▶ You don't have to take my word for this, though. Let's see it in action.

Let's create a “population” of 100,000 random numbers from a normal distribution with mean 0 and standard deviation 1

```
# note that rnorm is a "random" number generator, so we need to set a seed to make sure we get the same results each time  
set.seed(1304697)  
# NOTE: just set the seed once, at the top of your script. Then run everything and you will get reproducible results  
population <- rnorm(100000, mean = 0, sd = 1)  
  
mean(population)  
  
## [1] -0.003674513
```

- ▶ So we know the true mean is  $\mu = -0.0036745$
- ▶ Let's see what happens as we take a sample of size 10:

```
sample <- population[sample(1:length(population), 10, replace = FALSE)]
```

```
mean(sample)
```

```
## [1] -0.07407166
```

- Let's see what happens as we take a sample of size 100:

```
sample <- population[sample(1:length(population), 100, replace = FALSE)]
```

```
mean(sample)
```


```
## [1] -0.01106315
```



## Suppose we did this a bunch of times...

- ▶ Let's say we took a sample of size 10, 100 times
  - ▶ We would get 100 different sample means
  - ▶ We can calculate how “far” each sample mean is from the true mean
  - ▶ We can do the same thing for a bunch of different sample sizes
- ▶ Let's then plot the average mean squared error (MSE) for each sample size
  - ▶ The MSE is the average squared difference between the sample mean and the true mean
  - ▶ In other words, it is the average of  $(\bar{x} - \mu)^2$

## The results

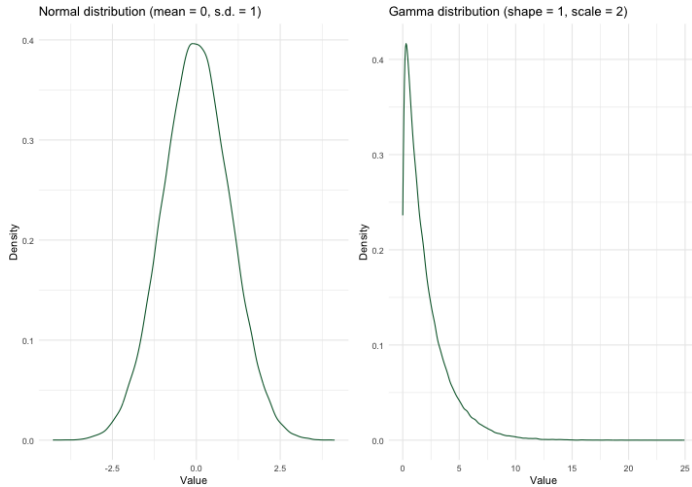


- ▶ This doesn't just work with normal distributions
- ▶ Here's an example with a more skewed distribution
- ▶ Let's create a "population" of 100,000 random numbers from a gamma distribution with a shape ( $k$ ) of 1 and a scale ( $\theta$ ) of 2


```
population2 <- rgamma(100000, shape = 1, scale = 2)  
mean(population2)
```

```
## [1] 2.005826
```

# The two populations



## The results



## The takeaway from the LLN

- ▶ The main takeaway from the LLN is that as the sample size increases, the sample mean will converge to the population mean
  - ▶ In other words, as  $n \rightarrow \infty$ ,  $\bar{x} \rightarrow \mu$
- ▶ This holds for many different distributions!
  - ▶ Something we will return to is that this also holds for the distribution of  $\hat{\beta}$


- ▶ The LLN is about the **mean**
  - ▶ The Central Limit Theorem (CLT) is about the **distribution**
- ▶ The CLT says that as the sample size increases, the distribution of the sample mean will converge to a normal distribution
  - ▶ In other words, as  $n \rightarrow \infty$ ,  $\bar{x} \sim N(\mu, \sigma^2/n)$
  - ▶ where  $\mu$  is the true population mean and  $\sigma^2$  is the true population variance

## The CLT, empirically

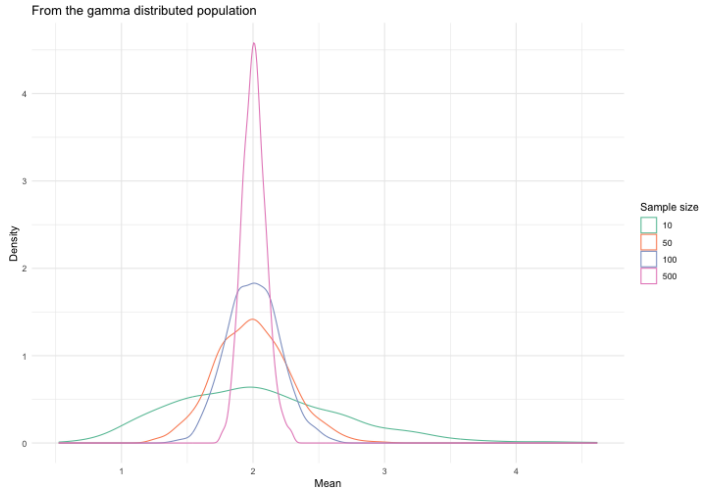
- ▶ We can do something similar to what we did with the LLN
  - ▶ Instead of looking at means, though we will look at *distributions of the mean*
  - ▶ In other words, we will density functions of the sample means
- ▶ I am going to take a sample and find the mean
  - ▶ Then I'm going to do it again
  - ▶ And again
  - ▶ 1,000 times
- ▶ Then I will plot the density of the sample means, for four separate sample sizes (10, 50, 100, and 500)



# The CLT with four sample sizes: 10, 50, 100, and 500 and 1,000 replications



# The CLT with four sample sizes: 10, 50, 100, and 500 and 1,000 replications



## The CLT lets us quantify uncertainty

- ▶ The key thing about the CLT is that the math behind us *lets us quantify the uncertainty!*
- ▶ For sample means, the CLT says that the standard error of the mean is:
  - ▶  $\sigma_{\bar{x}} = \frac{\sigma}{\sqrt{n}}$ 
    - ▶ where  $\sigma$  is the population standard deviation and  $n$  is the sample size
- ▶ Of course, we don't know  $\sigma$ , but we can estimate it!

## The CLT lets us quantify uncertainty

- ▶ We can estimate  $\sigma$  with  $s$ 
  - ▶  $s$  is the sample standard deviation
  - ▶ So we can estimate the standard error of the mean as:
    - ▶  $SE_{\bar{x}} = \frac{s}{\sqrt{n}}$
- ▶ This “standard error of the mean” is the standard deviation of the distribution of the sample mean
  - ▶ That distribution is called the sampling distribution of the mean

- ▶ Let's go back to our (population) regression equation:

$$\mathbf{Y} = \mathbf{X}\beta + \epsilon \quad (8)$$

- ▶ What drives uncertainty here?
  - ▶ The **error term**,  $\epsilon$
  - ▶ Note that an error term is a *population* parameter
  - ▶ The sample analog is the **residual**

- ▶ We can see that the error term is responsible for uncertainty by thinking about the deviation of estimated regression coefficients ( $\hat{\beta}$ ) from the true regression coefficients ( $\beta$ )

$$\mathbf{Y} = \beta\mathbf{X} + \epsilon \qquad \hat{\beta} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Y}$$

$$\hat{\beta} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'(\beta\mathbf{X} + \epsilon)$$

$$\hat{\beta} = \beta + (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\epsilon \qquad (9)$$

- ▶  $\beta$  is fixed, but  $\epsilon$  is random
  - ▶ So  $\hat{\beta}$  is random, too
  - ▶ That randomness is driven by the error term,  $\epsilon$

- ▶ Variance is about how much a random variable varies from its mean
  - ▶ It's a second-order moment: it's about how much a random variable varies from its mean (true value!)
  - ▶ Note how we use expectations here, since we are talking about population parameters

$$\text{Var}(\hat{\beta}|\mathbf{X}) = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbb{E}(\epsilon'\epsilon|\mathbf{X})\mathbf{X}(\mathbf{X}'\mathbf{X})^{-1} \quad (10)$$

- ▶ Well that doesn't look very good, does it?
  - ▶ We can simplify this. . . with an assumption on  $\mathbb{E}(\epsilon'\epsilon|\mathbf{X})$ , which we can denote  $\Omega$

- If we are willing to assume structure on the error term, we can simplify the variance

Assumption: Homoskedasticity

$\Omega = \sigma^2 \mathbf{I}$ , where  $\mathbf{I}$  is the identity matrix (ones on the diagonal and zeros elsewhere).

This simplifies the variance to:

$$\text{Var}(\hat{\beta}_{\text{homoskedasticity}}) = \sigma^2 (\mathbf{X}'\mathbf{X})^{-1} \quad (11)$$



- ▶ What, exactly, is homoskedasticity?
- ▶ It means that the variance of the error term is constant
  - ▶ In other words, the variance of the error term does not depend on the covariates
  - ▶ This is a very strong assumption and is often violated in practice

- ▶ Even though it may not be reasonable, let's assume homoskedasticity
  - ▶ Then, the variance of  $\hat{\beta}_{homoskedasticity}$  is:

$$\sigma^2(\mathbf{X}'\mathbf{X})^{-1} \quad (12)$$

- ▶ There's a problem, though: we don't know  $\sigma^2$ 
  - ▶ We can estimate it, just like with the mean
  - ▶ We can estimate it with the residual sum of squares (RSS)
    - ▶ The RSS is the sum of the squared *residuals*

## Estimating the variance under homoskedasticity

- ▶ We can estimate  $\sigma^2$  as  $\frac{RSS}{n-k-1}$ , where  $n$  is the sample size and  $k$  is the number of covariates
  - ▶ The extra 1 is due to the intercept
- ▶ We estimate the variance as:

$$Var(\hat{\beta}_{homoskedasticity}) = \hat{\sigma}^2(\mathbf{X}'\mathbf{X})^{-1} \quad (13)$$

$$= \frac{RSS}{n-k-1}(\mathbf{X}'\mathbf{X})^{-1} \quad (14)$$

$$= \frac{\sum_1^n (y_i - \hat{y}_i)^2}{n-k-1}(\mathbf{X}'\mathbf{X})^{-1} \quad (15)$$

## What does this equation tell us?

$$\text{Var}(\hat{\beta}_{\text{homoskedasticity}}) = \frac{\sum_1^n (y_i - \hat{y}_i)^2}{n - k - 1} (\mathbf{X}'\mathbf{X})^{-1} \quad (16)$$

- ▶ There are a couple things to point out here
  - ▶ First, the variance of  $\hat{\beta}$  is a function of the variance of the error term
    - ▶ The more unexplained variation in the outcome, the larger the variance of our estimated coefficients
  - ▶ Second, the variance of  $\hat{\beta}$  is also a function of the variance of the covariates
    - ▶ The more variation in the covariates, the smaller the variance of our estimated coefficients

## Why do we care about variance?

- ▶ We care about variance because it allows us to do hypothesis testing
  - ▶ We can test whether or not a coefficient is different from zero (or any other value)
  - ▶ We can test whether or not two coefficients are different from each other
- ▶ We only have a *sample*, which means we cannot say anything with certainty
  - ▶ We can only say something with a certain degree of confidence
- ▶ Taking into account variance allows us to say something about *the population from which the sample is drawn*
  - ▶ We generally don't care about the sample itself, only what it tells us about something larger

- ▶ Suppose we have an estimated coefficient,  $\hat{\beta}$
- ▶ Let's create a simulation exercise with a “population”

```
# some variable, x, randomly distributed between 0 and 5
x <- runif(100000, min = 0, max = 5)
# y is a function of x, plus some random error
y <- 3*x + rnorm(100000, mean = 0, sd = 1)
# put it into a data frame
df <- as_tibble(cbind(y, x))

# what is the "true" value of beta in our population of 100,000 people?
lm(y ~ x, data = df)
beta <- coef(lm(y ~ x, data = df))[2]
beta
```

## Confidence intervals

```
# what is the "true" value of beta in our population of 100,000 people?  
# I'm ignoring uncertainty for now  
lm(y ~ x, data = df)
```

```
##  
## Call:  
## lm(formula = y ~ x, data = df)  
##  
## Coefficients:  
## (Intercept)          x  
## -0.006007      3.003164  
beta <- coef(lm(y ~ x, data = df))[2]  
beta
```

```
##          x  
## 3.003164
```

Now let's take a sample of just 50. In a smaller sample, the coefficient can be quite different sometimes.


```
sample <- df[sample(1:nrow(df), 50, replace = FALSE),]  
summary(lm(y ~ x, data = sample))
```

```
##  
## Call:  
## lm(formula = y ~ x, data = sample)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -1.95124 -0.67324  0.05356  0.61053  2.52386   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)   0.3640     0.3214   1.133   0.263      
## x             2.9222     0.1045  27.971 <2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 1.026 on 48 degrees of freedom  
## Multiple R-squared:  0.9422, Adjusted R-squared:  0.941  
## F-statistic: 782.4 on 1 and 48 DF,  p-value: < 2.2e-16
```



- ▶ We want to put a confidence interval around our estimate of  $\hat{\beta}$ 
  - ▶ This will give us some idea of what the “true” value is (we know it here, but we generally do not!)
- ▶ You hopefully remember two separate distributions from earlier classes:
  - ▶ z distribution
  - ▶ t distribution
- ▶ In theory, you should always use a t distribution if *you do not know the true population  $\sigma^2$*  (which we usually don't)
  - ▶ We will always use t distributions

## Different distributions




- ▶ I'll assume you all remember critical values from earlier metrics classes
  - ▶ If not, ask now!
  - ▶ Confidence level (e.g. 0.05 for 95% confidence):  $\alpha$

- ▶ I'll assume you all remember critical values from earlier metrics classes
  - ▶ If not, ask now!
  - ▶ Confidence level (e.g. 0.05 for 95% confidence):  $\alpha$
- ▶ We construct our CI as:

$$CI = \left( \hat{\beta} - t_{n-k-1}^{1-\frac{\alpha}{2}} \times \sqrt{Var(\hat{\beta})}, \hat{\beta} + t_{n-k-1}^{1-\frac{\alpha}{2}} \times \sqrt{Var(\hat{\beta})} \right) \quad (17)$$


- ▶ At smaller sample sizes, the t distribution has fatter tails
  - ▶ This means that we need to be more uncertain about the true value of  $\beta$  and the critical value can be quite different

## Critical values for the different distributions



- ▶ Let's go back to our example with the sample of 50
- ▶ What, exactly, does the confidence interval represent?
  
- ▶ Suppose we took a bunch of samples of size 50, let's say 100 of them
  - ▶ We would get a bunch of different estimates of  $\hat{\beta}$
  - ▶ Let's plot them along with the confidence intervals

Degrees of freedom:  $n - k - 1 = 50 - 1 - 1 = 48$



```
## [1] "Coverage rate is defined as the proportion of samples for which the CI contains the true value: 0.94"
```

- ▶ Confidence intervals are a frequentist idea
  - ▶ They are about the probability of the true value being in the interval
- ▶ We usually can't know for sure whether or not the true value is in the interval
  - ▶ The probability is only about “averages” over many samples
- ▶ We can change our confidence level:
  - ▶  $\alpha = 0.01$  (confidence level of 99%) gives us a wider interval
  - ▶  $\alpha = 0.10$  (confidence level of 90%) gives us a narrower interval



- ▶ We can also explicitly test whether or not a coefficient is different from some value
  - ▶ What is the most common value we are interested in?

- ▶ We can also explicitly test whether or not a coefficient is different from some value
  - ▶ What is the most common value we are interested in?
  - ▶ Zero!
- ▶ We can do this with a t-test
  - ▶ Null hypothesis ( $H_0$ ): what we assume is true
  - ▶ Alternative hypothesis ( $H_1$ ): not the null
- ▶ Common example with a regression coefficient:
  - ▶  $H_0 : \beta = 0$
  - ▶  $H_1 : \beta \neq 0$
- ▶ NOTE: We are testing for the POPULATION parameter, not the sample statistics

- t-test:

$$\hat{t} = \frac{\hat{\beta} - H_0}{SE(\hat{\beta})}, \quad (18)$$

where  $SE(\hat{\beta})$  is the standard error of  $\hat{\beta}$ , or  $\sqrt{Var(\hat{\beta})}$ . I use  $\hat{t}$  here to underline that this comes from our sample statistics.

- We can then compare this to a critical value
  - If  $|\hat{t}| > t_{n-k-1}^{1-\frac{\alpha}{2}}$ , then we reject the null
  - If  $|\hat{t}| < t_{n-k-1}^{1-\frac{\alpha}{2}}$ , then we fail to reject the null

## Type 1 and type 2 errors

- ▶ Type 1 error: we reject the null when it is true
  - ▶  $\alpha$  is the probability of a type 1 error
- ▶ Type 2 error: we fail to reject the null when it is false
  - ▶ We do not know the probability of a type 2 error

## Interpreting the output

```
# using feols (from fixest package) instead of lm now
feols(agemployment ~ log(gnipc) + lfp + log(population) + urbanpop, data = data)
```

```
## OLS estimation, Dep. Var.: agemployment
## Observations: 1,420
## Standard-errors: IID
##
```

|                    | Estimate   | Std. Error | t value   | Pr(> t )      |
|--------------------|------------|------------|-----------|---------------|
| ## (Intercept)     | 105.201328 | 2.731881   | 38.50875  | < 2.2e-16 *** |
| ## log(gnipc)      | -10.112581 | 0.239763   | -42.17745 | < 2.2e-16 *** |
| ## lfp             | 0.241674   | 0.015230   | 15.86875  | < 2.2e-16 *** |
| ## log(population) | 0.376890   | 0.129339   | 2.91398   | 0.0036246 **  |
| ## urbanpop        | -0.208143  | 0.016205   | -12.84437 | < 2.2e-16 *** |

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 7.33605   Adj. R2: 0.834605
```

## Something really nice about fixest (feols)

```
# using feols instead of lm now
reg1 <- feols(agemployment ~ log(gnipc) + lfp + log(population) + urbanpop, data = data)
reg2 <- feols(serviceemployment ~ log(gnipc) + lfp + log(population) + urbanpop, data = data)
etable(reg1, reg2,
  se.below = TRUE, se.row = FALSE,
  digits = 3,
  digits.stats = 3,
  fitstat = c("r2", "ar2", "f", "n")
)
```

```
##                reg1                reg2
## Dependent Var.: agemployment serviceemployment
##
## Constant          105.2***          -7.05**
##                   (2.73)            (2.18)
## log(gnipc)        -10.1***           7.35***
##                   (0.240)           (0.191)
## lfp                0.242***          -0.048***
##                   (0.015)           (0.012)
## log(population)    0.377**           -0.873***
##                   (0.129)           (0.103)
## urbanpop          -0.208***           0.267***
##                   (0.016)           (0.013)
## -----
## R2                 0.835             0.849
## Adj. R2            0.835             0.849
## F-test             1,791.1           1,995.2
## Observations       1,420             1,420
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## Some important statistics


- ▶  $R^2$ : the proportion of variation in the outcome explained by the model
  - ▶  $R^2$  is always between 0 and 1
  - ▶  $R^2$  is a measure of *fit*, but you don't really want to use it to judge how “good” your model is
  - ▶ Adding more variables to your model can NEVER decrease  $R^2$
- ▶ Adjusted  $R^2$  ( $\bar{R}^2$ ): penalizes  $R^2$  for adding more variables
  - ▶ Adjusted  $R^2$  is always between 0 and 1
  - ▶ Adjusted  $R^2$  is a measure of *fit*, but you don't really want to use it to judge how “good” your model is
  - ▶ Adding more variables to your model can decrease adjusted  $R^2$

$$\bar{R}^2 = 1 - (1 - R^2) \frac{n - 1}{n - k - 1} \quad (19)$$

- ▶ The F distribution is important in linear regression
  - ▶ It is the ratio of two chi-squared distributions, so it has two different degrees of freedom
  - ▶ We sometimes refer to these as the “numerator” and “denominator” degrees of freedom
  - ▶ An F statistic is always positive
- ▶ In the previous regression, the F at the bottom is a test of whether or not all of the coefficients are equal to zero
  - ▶ This is a joint test of the coefficients
  - ▶ The null hypothesis is that *all* of the coefficients are equal to zero
  - ▶ The alternative hypothesis is that *at least one* of the coefficients is not equal to zero
- ▶ In practice, we usually just use the p-value because of the fact that critical values of F distributions can be quite different depending on the d.f.



## F distributions with different degrees of freedom



- ▶ t-tests are the workhorse for inference on a *single* coefficient
- ▶ But what if we want to test multiple coefficients at once?
  - ▶ F-tests
- ▶ Consider the following regression:

$$y_i = \beta_0 + \beta_1 x_{1,i} + \beta_2 x_{2,i} + \beta_3 x_{3,i} + \epsilon_i \quad (20)$$

- ▶ How might we think about testing whether  $\beta_1 = \beta_2 = 0$ ?

$$y_i = \beta_0 + \beta_1 x_{1,i} + \beta_2 x_{2,i} + \beta_3 x_{3,i} + \epsilon_i \quad (21)$$

- If we assume  $\beta_1 = \beta_2 = 0$ , that is the same as estimating:

$$y_i = \tilde{\beta}_0 + \tilde{\beta}_1 x_{1,i} + \tilde{\epsilon}_i \quad (22)$$

- Intuitively, we want to know whether the first specification is “better” than the second. We can calculate the estimated variance of the residuals in each specification:

$$\hat{\sigma}^2 = \sum_1^n (y_i - \hat{y}_i)^2 \quad (23)$$

$$\tilde{\sigma}^2 = \sum_1^n (y_i - \tilde{y}_i)^2 \quad (24)$$

$$\hat{\sigma}^2 = \sum_1^n (y_i - \hat{y}_i)^2 \quad (25)$$

$$\tilde{\sigma}^2 = \sum_1^n (y_i - \tilde{y}_i)^2 \quad (26)$$

- We can then calculate the F statistic:

$$F = \frac{(\tilde{\sigma}^2 - \hat{\sigma}^2)/q}{\hat{\sigma}^2/(n - k - 1)}, \quad (27)$$

where  $q$  is the number of restrictions we are testing. In this case,  $q = 2$  (two restricted coefficients).

$$F = \frac{(\tilde{\sigma}^2 - \hat{\sigma}^2)/q}{\hat{\sigma}^2/(n - k - 1)} \quad (28)$$

- ▶ The numerator  $((\tilde{\sigma}^2 - \hat{\sigma}^2)/q)$  is distributed as  $\chi^2$  with  $q$  degrees of freedom
- ▶ The denominator  $(\hat{\sigma}^2/(n - k - 1))$  is distributed as  $\chi^2$  with  $n - k - 1$  degrees of freedom
- ▶ Their ratio is distributed as F with  $q$  and  $n - k - 1$  degrees of freedom:

$$F = \frac{(\tilde{\sigma}^2 - \hat{\sigma}^2)/q}{\hat{\sigma}^2/(n - k - 1)} \sim F(q, n - k - 1) \quad (29)$$

- ▶ Now you see why they are sometimes referred to as “numerator” and “denominator” degrees of freedom!

## F-test by hand - note the relationship with the t-test for a single coefficient!

```
# full specification
full <- feols(agemployment ~ log(gnipc) + lfp + log(population) + urbanpop, data = data)
# restricted
restricted <- feols(agemployment ~ log(gnipc) + lfp + urbanpop, data = data)
sigma_full <- sum((full$resid)^2) # sigma for full
sigma_restricted <- sum((restricted$resid)^2) # sigma for restricted
F <- ((sigma_restricted - sigma_full)/1)/(sigma_full/(full$nobs - 4 - 1)) # calculate F statistic
F
```

```
## [1] 8.49128
```

```
full
```

```
## OLS estimation, Dep. Var.: agemployment
## Observations: 1,420
## Standard-errors: IID
##
```

|                    | Estimate   | Std. Error | t value   | Pr(> t )      |
|--------------------|------------|------------|-----------|---------------|
| ## (Intercept)     | 105.201328 | 2.731881   | 38.50875  | < 2.2e-16 *** |
| ## log(gnipc)      | -10.112581 | 0.239763   | -42.17745 | < 2.2e-16 *** |
| ## lfp             | 0.241674   | 0.015230   | 15.86875  | < 2.2e-16 *** |
| ## log(population) | 0.376890   | 0.129339   | 2.91398   | 0.0036246 **  |
| ## urbanpop        | -0.208143  | 0.016205   | -12.84437 | < 2.2e-16 *** |

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 7.33605  Adj. R2: 0.834605
```

- For a *single* coefficient,  $F = t^2$

- ▶ In the previous examples, we've been assuming *homoskedasticity*

$$\mathbb{E}(\epsilon'\epsilon|\mathbf{X}) = \mathbf{\Omega} = \sigma^2\mathbf{I} \quad (30)$$

- ▶ In practice, this assumption is unlikely to be true
  - ▶ You can test for it, but it's not really worth it. Just assume it's not true.
  - ▶ In other words, assume  $\mathbf{\Omega} \neq \sigma^2\mathbf{I}$
- ▶ **Homoskedasticity** can take many forms
  - ▶ It simply means that the variance is not constant across observations
  - ▶ Let's talk about a general estimator under heteroskedasticity and another one under clustering

- ▶ There is a broad class of estimators that are called “sandwich estimators” because of their form
- ▶ The baseline *heteroskedasticity-consistent* standard errors are:

$$\text{Var}(\hat{\beta}) = (\mathbf{X}'\mathbf{X})^{-1} \left( \sum_i^n X_i X_i' \sigma_i^2 \right) (\mathbf{X}'\mathbf{X})^{-1} \quad (31)$$

- ▶ Of course we don't know  $\sigma_i^2$ , so we calculate a feasible estimate using the residuals:

$$\text{Var}(\hat{\beta})^{HC0} = (\mathbf{X}'\mathbf{X})^{-1} \left( \sum_i^n X_i X_i' \hat{e}_i^2 \right) (\mathbf{X}'\mathbf{X})^{-1} \quad (32)$$

- ▶ This is sometimes referred to as *HC0* standard errors, or **Eicker-White** standard errors



- ▶ But there's a problem: it turns out that  $\sum_i \hat{e}_i^2$  is a biased estimator for  $\sigma_i^2$ !
  - ▶ It's biased towards zero ("attenuated") so the standard errors are a bit too small
- ▶ The next step is to use a scaling factor to "correct" the standard errors
  - ▶ This is sometimes referred to as *HC1* standard errors

$$Var(\hat{\beta})^{HC1} = \frac{n}{n - k - 1} (\mathbf{X}'\mathbf{X})^{-1} \left( \sum_i^n X_i X_i' \hat{e}_i^2 \right) (\mathbf{X}'\mathbf{X})^{-1} \quad (33)$$

- ▶ This is a bit ad hoc, but is generally preferred (Hanson, 2022).
  - ▶ **This is currently the default in Stata (with the option *robust*)**

- ▶ To confuse you more, there are two additional alternatives:

$$\text{Var}(\hat{\beta})^{HC2} = (\mathbf{X}'\mathbf{X})^{-1} \left( \sum_i^n (1 - h_{ii})^{-1} X_i X_i' \hat{e}_i^2 \right) (\mathbf{X}'\mathbf{X})^{-1} \quad (34)$$

$$\text{Var}(\hat{\beta})^{HC3} = (\mathbf{X}'\mathbf{X})^{-1} \left( \sum_i^n (1 - h_{ii})^{-2} X_i X_i' \hat{e}_i^2 \right) (\mathbf{X}'\mathbf{X})^{-1} \quad (35)$$

- ▶  $h_{ii}$  has to do with a diagonal element of an orthogonal projection matrix. We're going to ignore this for now.
- ▶ These are sometimes referred to as *HC2* and *HC3* standard errors
- ▶ When I refer to HC standard errors, I will be talking about *HC1* unless I say otherwise

- ▶ One of the most common types of heteroskedasticity is clustering
  - ▶ This is when the error term is correlated within groups
  - ▶ For example, if we have data on students within schools, the error term for students in the same school is likely to be correlated
- ▶ Let  $K$  denote the number of clusters and  $k$  a specific cluster
  - ▶ We can then calculate the cluster-robust variance estimator as:

$$Var(\hat{\beta})^{cluster} = (\mathbf{X}'\mathbf{X})^{-1} \sum_{k=1}^K \left( \sum_{i=1}^{n_k} X_{ik} \hat{e}_{ik} \right) \left( \sum_{l=1}^{n_k} X_{lk} \hat{e}_{lk} \right)' (\mathbf{X}'\mathbf{X})^{-1} \quad (36)$$

- ▶ The basic intuition is that the more correlated the observations in a cluster is, the larger the variance of the estimated coefficients
  - ▶ You can actually show this with some algebra and some assumptions
  - ▶ This *intra-cluster correlation* is often referred to as  $\rho$

- ▶ Suppose you have no clustering at all. Then your effective sample size is  $n$ .
- ▶ Suppose all of your observations are perfectly correlated with the cluster ( $\rho = 1$ ). Then your effective sample size is just the number of clusters.
- ▶ Of course, neither assumption is likely to ever be true. We are usually somewhere in between, but hopefully this helps with intuition.
- ▶ All of these alternative standard error (variance) calculations only affect the standard errors. **Coefficients do not change.**

## Clustering and standard errors

```
reg1 <- feols(agemployment ~ log(gnipc) + lfp + log(population) + urbanpop, data = data) # assume iid
reg2 <- feols(agemployment ~ log(gnipc) + lfp + log(population) + urbanpop, data = data, cluster = "country") # clustering
etable(reg1, reg2,
        se.below = TRUE, se.row = FALSE,
        digits = 3, digits.stats = 3,
        fitstat = c("r2", "ar2", "f", "n")
)
```

```
##                reg1                reg2
## Dependent Var.: agemployment agemployment
##
## Constant          105.2***          105.2***
##                   (2.73)           (9.40)
## log(gnipc)        -10.1***          -10.1***
##                   (0.240)          (0.713)
## lfp                0.242***          0.242***
##                   (0.015)          (0.041)
## log(population)    0.377**           0.377
##                   (0.129)          (0.462)
## urbanpop          -0.208***          -0.208***
##                   (0.016)          (0.047)
## -----
## R2                 0.835            0.835
## Adj. R2            0.835            0.835
## F-test             1,791.1          178.5
## Observations       1,420            1,420
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## What level do we cluster at?

- ▶ This is a hard question to answer in many cases
  - ▶ With panel data, we usually cluster at the unit level (e.g. individuals)
  - ▶ In the previous example, this was the country level
- ▶ Always cluster at the level of randomization
  - ▶ If the data is collected using a complex survey design, we cluster at the the “enumeration area” level
- ▶ **We do not know the “correct” level at which to cluster except in rare situations.**

- ▶ What if the size of clusters is quite different?
  - ▶ Variance is calculated as the sum of errors within clusters, so different cluster sizes means differences in variance
  - ▶ This is basically heteroskedasticity
  - ▶ The key takeaway from this is that smaller samples may show bias (finite-sample bias)
  - ▶ Which leads to...
- ▶ What if  $K$  is small, i.e. we have relatively few clusters?
  - ▶ One way to think about it is that our effective sample size is closer to  $K$  than  $n$
  - ▶ Unless we are willing to assume normality of the error term (which we already know we can't), small samples can be problematic
  - ▶ What to do? We will return to this in the next section on bootstrapping.

Have you seen these before?

1. **Linearity:** the relationship between the outcome and covariates is linear
2. **Full rank:** the matrix  $\mathbf{X}$  has full rank - Another way to say this: no perfect collinearity
3. **Exogeneity:** the error term is uncorrelated with the covariates
4. **Homoskedasticity:** the variance of the error term is constant
5. **Normality:** the error term is normally distributed



1. **Linearity:** the relationship between the outcome and covariates is linear - Regardless of the “true” relationship, OLS will always give you the best linear approximation

2. **Full rank:** the matrix  $\mathbf{X}$  has full rank, or no perfect collinearity - This is required to estimate the regression. - If there is perfect collinearity, then the matrix  $\mathbf{X}'\mathbf{X}$  is not invertible - Stata and R will automatically drop variables to make this happen

3. **Exogeneity:** the error term is uncorrelated with the covariates - Ignore this for now. We'll come back to it.

4. **Homoskedasticity:** the variance of the error term is constant - Obviously not required! We've already talked about this. There are corrections.

5. **Normality:** the error term is normally distributed - This is not required for OLS! - The only time this is important is with small samples. - With small samples, we need to assume normality to do hypothesis testing. - With large samples, we can use the central limit theorem to get around non-normality. - A note: the more “non-normal” the error term is, the larger the sample size we need for appropriate inferences.

## Bootstrapping (resampling)

- ▶ It turns out that there is an alternative to calculating standard errors using formulas
  - ▶ It's called bootstrapping
  - ▶ It's a resampling method
    - ▶ I'll abuse terminology and will generally refer to all of these as “bootstrapping”
- ▶ Bootstrapping originated with Efron (1970)
  - ▶ Lots of work since then
  - ▶ There are many different types of bootstrapping

- ▶ Bootstrapping is about sampling from your data
  - ▶ Something you might not be used to: it is sampling *with replacement*
- ▶ Imagine a vector with 10 elements:  $\{1, 2, 3, 4, 5, 6, 7, 8, 9, 10\}$ 
  - ▶ Imagine we sample a randomly selected element. Let's say we get 5.
  - ▶ Now, we want to take another randomly selected element. The question: do we put the 5 back in or not?
  - ▶ With replacement: put the 5 back in
  - ▶ Without replacement: don't put the 5 back in
- ▶ Bootstrapping is *with replacement*
  - ▶ Put the 5 back in!

## Resampling with replacement

```
## # A tibble: 10 x 2
##   country year
##   <chr>   <dbl>
## 1 AUT     2003
## 2 BGD     2003
## 3 BEL     2003
## 4 BEN     2003
## 5 BRA     2003
## 6 BGR     2003
## 7 KHM     2003
## 8 CAN     2003
## 9 CHL     2003
## 10 COL    2003
```



## Resampling 10 observations *with replacement*

- The key is that we sample rows – i.e. observations. We do not sample columns. We take an observation and all its values.

```
## # A tibble: 10 x 2
##   country year
##   <chr>   <dbl>
## 1 CHL     2003
## 2 COL     2003
## 3 KHM     2003
## 4 BEN     2003
## 5 AUT     2003
## 6 AUT     2003
## 7 BRA     2003
## 8 CAN     2003
## 9 BEN     2003
## 10 CAN    2003
```

- Suppose we are interested in the variance of a coefficient estimate,  $\hat{\beta}_{a_1}$ 
  - We can use bootstrapping to estimate the variance of  $\hat{\beta}_{a_1}$
  - Keeping it simple, here is our regression:

$$y_i = \beta_0 + \beta_1 x_{1,i} + \beta_2 x_{2,i} + \beta_3 x_{3,i} + \beta_4 x_{4,i} + \epsilon_i \quad (37)$$

```
feols(agemployment ~ log(gnipc) + lfp + log(population) + urbanpop, data = data)
```

```
## OLS estimation, Dep. Var.: agemployment
## Observations: 1,420
## Standard-errors: IID
##           Estimate Std. Error   t value Pr(>|t|)
## (Intercept)   105.201328    2.731881  38.50875 < 2.2e-16 ***
## log(gnipc)    -10.112581    0.239763 -42.17745 < 2.2e-16 ***
## lfp           0.241674    0.015230  15.86875 < 2.2e-16 ***
## log(population) 0.376890    0.129339   2.91398 0.0036246 **
## urbanpop     -0.208143    0.016205 -12.84437 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 7.33605   Adj. R2: 0.834605
```

- ▶ Imagine a single bootstrap sample, which we will denote with  $b$ 
  - ▶ We are going to call the estimate we are interested in  $\hat{\theta}_b$  (which in this case is a coefficient)
- ▶ Let's calculate the variance using our bootstrap as follows:

$$\text{Var}(\hat{\theta}) = \frac{1}{B-1} \sum_{b=1}^B (\hat{\theta}_b - \bar{\theta})^2, \quad (38)$$

where

$$\bar{\theta} = \frac{1}{B} \sum_{b=1}^B \hat{\theta}_b \quad (39)$$

- ▶ In words:
  - ▶ We take a bunch of bootstrap samples (more is better, but we'll do 1,000)
  - ▶ We estimate the regression in each bootstrap sample and save the coefficient we are interested in
  - ▶ After all 1,000 bootstrap samples, we calculate the variance as the mean (almost) of the squared differences between each bootstrap estimate and the mean of the bootstrap estimates

## Bootstrapping to estimate variance

```
# empty container to hold the estimates
bootvec <- c()
# 1000 samples
for (b in 1:1000){
  # take a sample with replacement (same size as data)
  sample <- data[sample(1:nrow(data), nrow(data), replace = TRUE),]
  # estimate regression
  reg <- feols(agemployment ~ log(gnipc) + lfp + log(population) + urbanpop, data = sample)
  # add coefficient to vector
  bootvec <- c(bootvec, reg$coefficients[2])
}
# find variance
var <- sum((bootvec - mean(bootvec))^2)/999
# estimated standard error is the square root
sqrt(var)
```

```
## [1] 0.2373331
```

- ▶ The estimated standard error is the square root of the variance, and it is quite close to the standard error we got from the regression
- ▶ In reality, there's no reason to do this for a single coefficient like this
  - ▶ But it's a good way to understand the intuition behind bootstrapping
- ▶ Where this becomes really powerful is where there is no closed-form solution for the variance
  - ▶ For example, what if we wanted to calculate  $\frac{\hat{\beta}_2}{\hat{\beta}_3}$ ?
  - ▶ To be clear, this is a non-sense example, but it's just to illustrate the point

## Our non-sense example

```
# empty container to hold the estimates
bootvec <- c()
# 1000 samples
for (b in 1:1000){
  # take a sample with replacement (same size as data)
  sample <- data[sample(1:nrow(data), nrow(data), replace = TRUE),]
  # estimate regression
  reg <- feols(employment ~ log(gnipc) + lfp + log(population) + urbanpop, data = sample)
  # add coefficient to vector
  bootvec <- c(bootvec, reg$coefficients[2]/reg$coefficients[3])
}
# find variance
var <- sum((bootvec - mean(bootvec))^2)/999
# what's the mean?
mean(bootvec)
```

```
## [1] -41.87397
```

```
# estimated standard error is the square root
sqrt(var)
```

```
## [1] 2.857473
```

- ▶ There's a problem with the previous example: we are assuming that the observations are independent
- ▶ What if we think there is clustering?
  - ▶ We can use a clustered bootstrap, where we randomly sample CLUSTERS, not observations
  - ▶ In this case, clusters are countries, so we will randomly sample countries



## Clustered bootstrap

```
# empty container to hold the estimates  
feols(agemployment ~ log(gnipc) + lfp + log(population) + urbanpop, data = data, cluster = "country")
```

```
## OLS estimation, Dep. Var.: agemployment  
## Observations: 1,420  
## Standard-errors: Clustered (country)  
##  
##      Estimate Std. Error   t value   Pr(>|t|)  
## (Intercept)  105.201328   9.400891  11.190570 < 2.2e-16 ***  
## log(gnipc)    -10.112581   0.713338 -14.176414 < 2.2e-16 ***  
## lfp           0.241674    0.041227   5.861974 3.0986e-08 ***  
## log(population) 0.376890    0.462084   0.815631 4.1609e-01  
## urbanpop     -0.208143    0.047444  -4.387083 2.2344e-05 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
## RMSE: 7.33605   Adj. R2: 0.834605
```

## Bootstrapping to estimate variance


```
# empty container to hold the estimates
bootvec <- c()
# get list of countries
countries <- unique(data$country)
# This takes longer to run, so just 500 as an example (try to use more)
for (b in 1:500){
  # take a sample with replacement (same size as data)
  samplecountries <- countries[sample(1:length(countries), length(countries), replace = TRUE)]
  # go through sample countries and get data, appending
  sample <- c()
  for (c in samplecountries){
    sample <- rbind(sample, data %>% filter(data$country==paste0(c)))
  }
  # estimate regression
  reg <- feols(agemployment ~ log(gnipc) + lfp + log(population) + urbanpop, data = sample)
  # add coefficient to vector
  bootvec <- c(bootvec, reg$coefficients[2])
}
```

```
## [1] "Standard error: 0.759"
```

- ▶ We can use our bootstrap vector to calculate confidence intervals
  - ▶ We can use the 2.5th and 97.5th percentiles of the bootstrap vector to get a 95% confidence interval:

$$CI = \left( \hat{\theta}_{\alpha/2}, \hat{\theta}_{\alpha/2} \right) \quad (40)$$

## Using the bootstrap vector for confidence intervals



```
## [1] "Mean: -10.2"
```

```
## [1] "Lower: -11.9"
```

```
## [1] "Upper: -8.75"
```

- ▶ This type of bootstrap is non-parametric
  - ▶ That is, we are not assuming anything about the distribution of the data
  - ▶ The previous kernel density plot wasn't normally distributed, it was skewed, etc.
  - ▶ This flexibility is one of the main advantages of bootstrapping
- ▶ The main downside to bootstrapping is computation
  - ▶ The previous example takes about 5 minutes to run on my computer with just 500 samples (and it's a good computer)
  - ▶ More elaborate estimation can take much longer

- ▶ Brummund and Merfeld (2022)
  - ▶ We want to estimate marginal revenue product of labor (MRPL)
- ▶ Here is the specification:

$$\begin{aligned} \ln R_{iht} = & \alpha_h + I(\text{Farm} = 1) \times \left( \sum_j \beta_j \ln \gamma_{jiht} + \frac{1}{2} \sum_j \sum_k \beta_{jk} \ln \gamma_{jiht} \ln \gamma_{kiht} + \delta C_{iht} \right. \\ & \left. + D_{dt} + \eta_m \right) + \sum_j \beta_j \ln \gamma_{jiht} + \frac{1}{2} \sum_j \sum_k \beta_{jk} \ln \gamma_{jiht} \ln \gamma_{kiht} + \delta C_{iht} + D_{dt} + \eta_m \\ & + I(\text{Farm} = 1) + \varepsilon_{iht} \end{aligned} \quad (41)$$

- This is what we need to estimate, using coefficients from the previous equation:

$$\frac{\partial R^f}{\partial L^f} = \frac{\hat{R}_{iht}^f}{L_{iht}^f} \left[ \beta_L^f + \beta_{LL}^f \log L_{iht}^f + \beta_{LA} \log A_{iht} + \beta_{LF} \log F_{iht} \right] \quad (42)$$

$$\frac{\partial R^{nf}}{\partial L^{nf}} = \frac{\hat{R}_{iht}^{nf}}{L_{iht}^{nf}} \left[ \beta_L^{nf} + \beta_{LL}^{nf} \log L_{iht}^{nf} + \beta_{LC} \log C_{iht} \right] \quad (43)$$

- It's not possible to get closed-form solutions for the variance of these
  - So what did we do? Bootstrapped!
  - Household panel data, so clustered bootstrap at the household level