

Microeconometrics (Causal Inference)

Weeks 5 and 6 - Differences-in-differences and synthetic control

Joshua D. Merfeld
KDI School of Public Policy and Management

2023-09-25

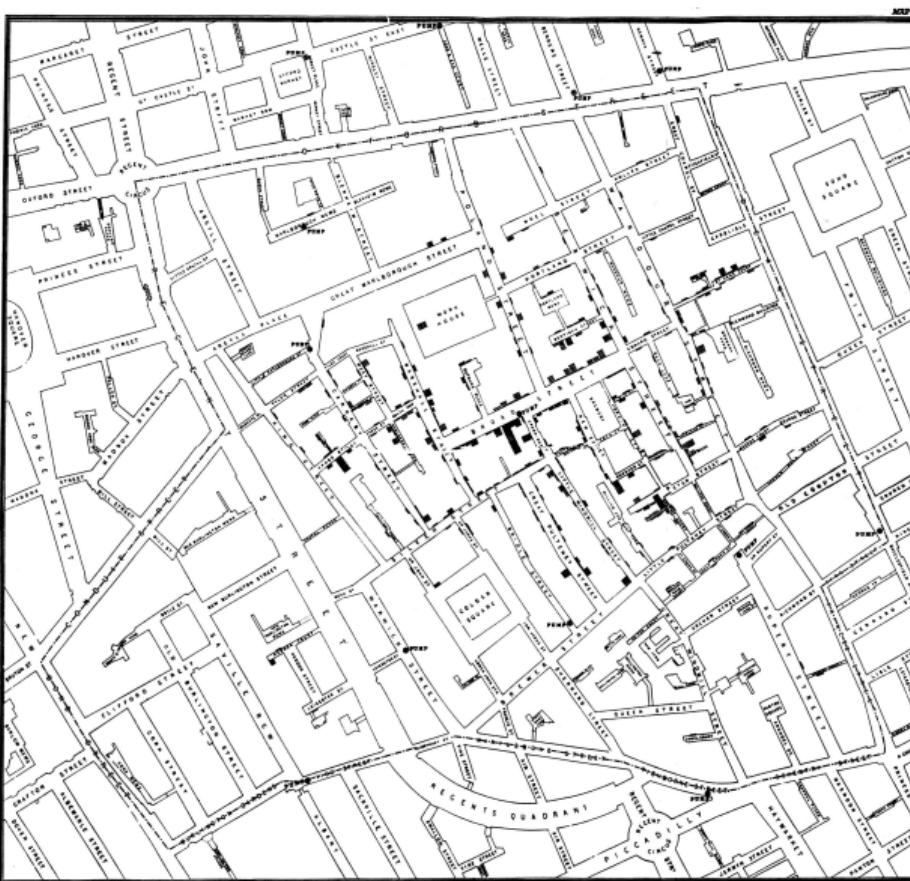
What are we doing today?

- ▶ Canonical differences-in-differences
 - ▶ Inference
 - ▶ Wild cluster bootstrap
- ▶ Fixed effects vs. random effects
- ▶ Bias in two-way fixed effects

Differences-in-differences... in the 1800s?



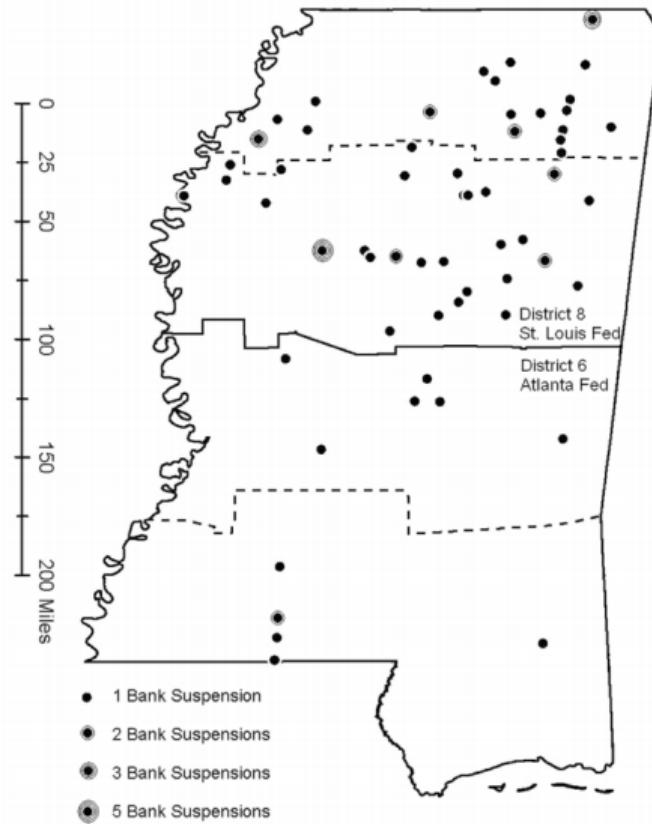
Differences-in-differences... in the 1800s? (from Cunningham's CI)



Differences-in-differences

- ▶ More commonly referred to as “DID” or “diff-in-diff”
 - ▶ Classic reference: Card and Krueger (1994)
- ▶ Most common method, likely because data requirements are least stringent
- ▶ Example in *Mostly Harmless*: offering credit to banks during the Great Depression (Richardson and Troost, 2009)
 - ▶ Set up: Two different federal reserve banks lent to neighborhood banks in Mississippi
 - ▶ Atlanta fed favored lending to banks in trouble
 - ▶ St. Louis fed favored the exact opposite

Richardson and Troost (2009) - Mississippi dividing line



Did the policy of extra lending save banks?

- ▶ Basic idea: compare what happened to Atlanta fed banks (southern Mississippi) with St. Louis fed banks (northern Mississippi)
- ▶ Could compare after lending, but what's the assumption here?

Did the policy of extra lending save banks?

- ▶ Basic idea: compare what happened to Atlanta fed banks (southern Mississippi) with St. Louis fed banks (northern Mississippi)
- ▶ Could compare after lending, but what's the assumption here?
- ▶ Assumption: same levels before intervention (very strict assumption)

In fact, pre-intervention levels are different!

TABLE 4
BANK SUSPENSIONS AND LIQUIDATIONS

Begin July 1	End June 30	All (1)	PERCENTAGE OF BANKS SUSPENDING			PERCENTAGE OF BANKS LIQUIDATING				
			Federal Reserve District			Federal Reserve District				
			6th	Atlanta	8th	St. Louis	All	(4)	6th	Atlanta
1929	to	1930	4.8	7.1	3.0	4.5	7.1	4.5	7.1	2.4
1930	to	1931	28.9	14.2	39.5	13.6	7.1	13.6	7.1	18.6
1931	to	1932	13.2	14.9	11.8	8.0	7.9	8.0	7.9	8.1
1932	to	1933	7.7	7.5	7.9	7.3	6.5	7.3	6.5	7.9
1933	to	1934	.9	.0	1.7	.9	.0	.9	.0	1.7
1929	to	1934 ^a	49.8	38.7	59.2	30.9	26.8	30.9	26.8	34.4

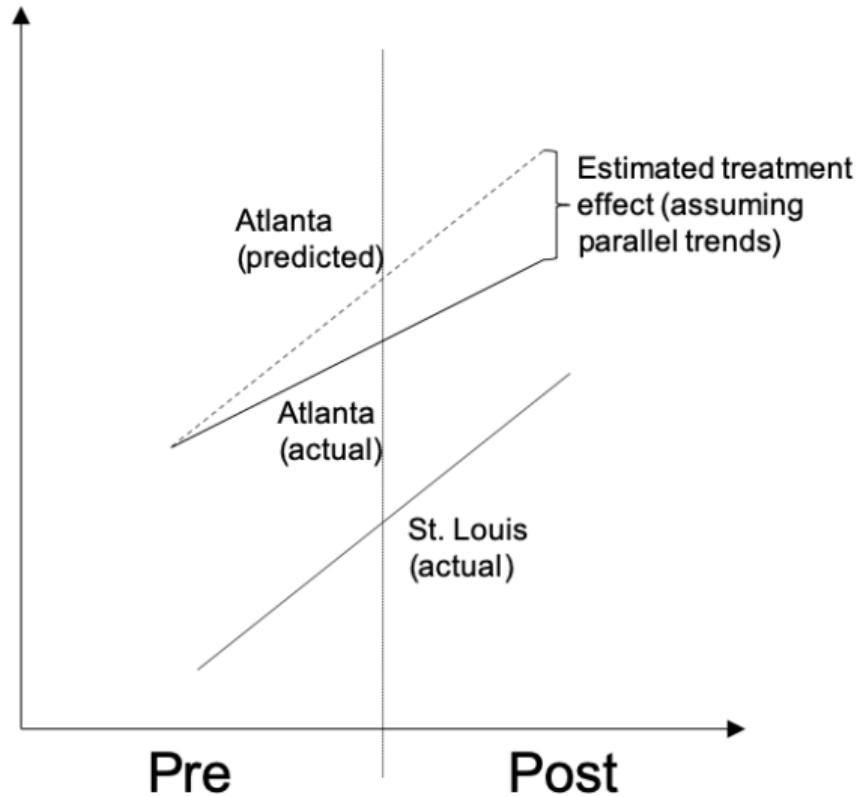
SOURCE.—*Rand McNally Bankers Directory* and National Archives and Records Administration Record Group 82. See Section II and Richardson (2006, 2007a, 2007b, 2008) for details.

^a The last row indicates the percentage of banks operating on July 1, 1929, that either suspended or liquidated by June 30, 1933.

Did the policy of extra lending save banks?

- ▶ Instead, compare *changes* from before to after treatment
- ▶ Assumption: parallel trends
- ▶ If valid, the fact the districts were different prior to the treatment isn't a problem

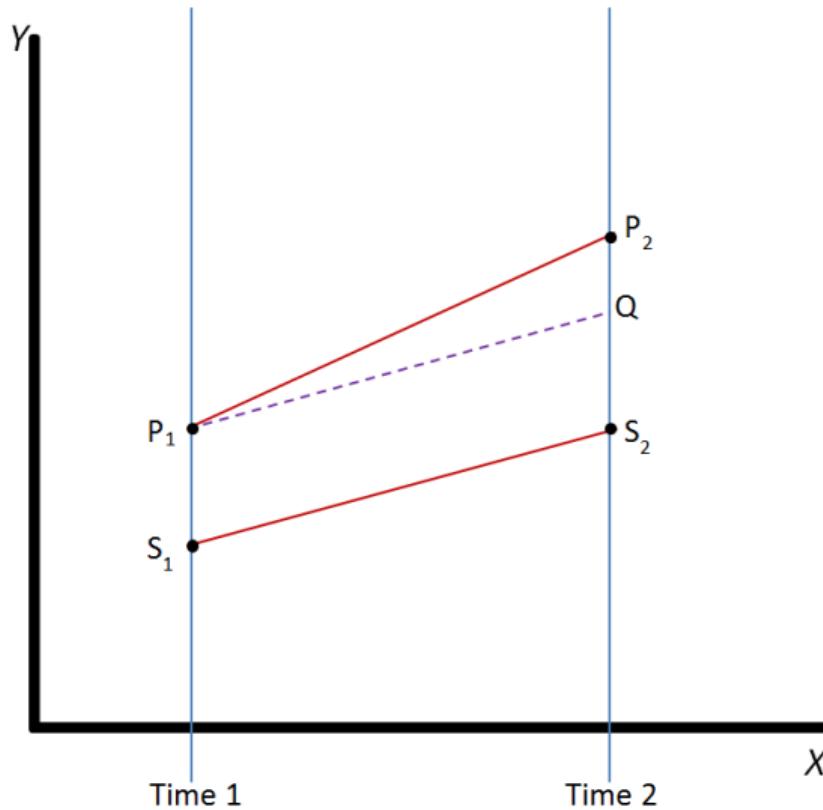
“Parallel trends”



Why is it “differences in differences”?

- ▶ Difference 1: St. Louis post minus St. Louis pre
- ▶ Difference 2: Atlanta post minus Atlanta pre
- ▶ Difference-in-differences: Difference 2 minus difference 1

“Differences in differences” graphically



Parallel trends assumption

- ▶ The key assumption in differences-in-differences is the parallel trends assumption
 - ▶ *If the treated group had not been treated, it would have changed by the same amount (“had the same trend”) as the comparison group.*
- ▶ This is a counterfactual assumption: We cannot explicitly test it
- ▶ What can we do instead?

Parallel trends assumption

- ▶ The key assumption in differences-in-differences is the parallel trends assumption
 - ▶ *If the treated group had not been treated, it would have changed by the same amount ("had the same trend") as the comparison group.*
- ▶ This is a counterfactual assumption: We cannot explicitly test it
- ▶ What can we do instead?
 - ▶ We can test trends *before* treatment
 - ▶ Or in the case of this article, *after* treatment!

Richardson and Troost (2009) - Testing the assumption

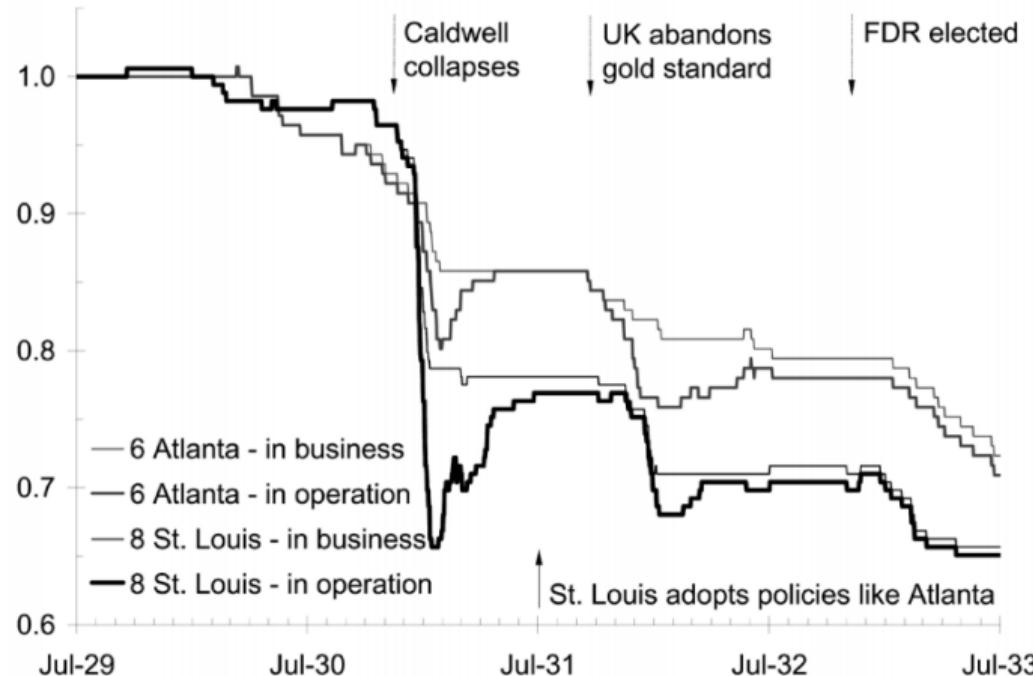


FIG. 3.—Percentage of banks in business and in operations in the 6th and 8th Federal Reserve Districts in Mississippi, July 1929 to June 1933. Source: See Section II.

Estimating diff-in-diff empirically

- ▶ Can be estimated in a straightforward regression:

$$Y_{it} = \beta_0 + \beta_1 TREAT_i + \beta_2 POST_t + \beta_3 (POST_t \times TREAT_i) + \varepsilon_{it} \quad (1)$$

- ▶ Can be estimated in a straightforward regression:

$$Y_{it} = \beta_0 + \beta_1 TREAT_i + \beta_2 POST_t + \beta_3 (POST_t \times TREAT_i) + \varepsilon_{it} \quad (1)$$

- ▶ β_0 : pre mean for the comparison group
- ▶ β_1 : difference in the pre mean between the treated and untreated group
- ▶ β_2 : difference in means between the pre and post period for the comparison group
- ▶ β_3 : difference-in-differences estimate
 - ▶ This is the difference in the change from pre to post for the treated group relative to the comparison group

Card and Krueger (1994) - Minimum wage and employment

```
ckdata <- read_csv("cardkruegerlong.csv")
head(ckdata)

## # A tibble: 6 x 15
##   sheet chain co_owned state southj centralj northj   pa1   pa2 shore calls
##   <dbl> <dbl>     <dbl> <dbl> <dbl>    <dbl> <dbl> <dbl> <dbl> <chr>
## 1    506     2       1     0     0      0     0     1     0     0 0
## 2     56     4       1     0     0      0     0     1     0     0 0
## 3     61     4       1     0     0      0     0     1     0     0 0
## 4     62     4       1     0     0      0     0     1     0     0 2
## 5    445     1       0     0     0      0     0     0     1     0 0
## 6    451     1       0     0     0      0     0     0     1     0 0
## # i 4 more variables: fulltime <chr>, parttime <chr>, managers <chr>,
## #   post <dbl>
# note that state = 1 for NJ and 0 for PA.
# also note that post = 1 for 1993 and 0 for 1992
# NJ is treated group, so state = 1 means treat = 1

# looks like fulltime is a character! let's try to make it numeric
ckdata$fulltime_num <- as.numeric(ckdata$fulltime)

## Warning: NAs introduced by coercion
ckdata$parttime_num <- as.numeric(ckdata$parttime)

## Warning: NAs introduced by coercion
```

Card and Krueger (1994) - Minimum wage and employment

```
# said there are NAs in the data, so let's see where they are
ckdata %>% filter(is.na(fulltime_num)) %>% select(fulltime, fulltime_num)

## # A tibble: 18 x 2
##   fulltime fulltime_num
##   <chr>     <dbl>
## 1 .          NA
## 2 .          NA
## 3 .          NA
## 4 .          NA
## 5 .          NA
## 6 .          NA
## 7 .          NA
## 8 .          NA
## 9 .          NA
## 10 .         NA
## 11 .         NA
## 12 .         NA
## 13 .         NA
## 14 .         NA
## 15 .         NA
## 16 .         NA
## 17 .         NA
## 18 .         NA

# ah, so they are .! those are missing values in Stata, so leave as missing.
```

Card and Krueger (1994) - Minimum wage and employment

```
reg1 <- feols(fulltime_num ~ state + post + state:post, data = ckdata, vcov = "HC1")
summary(reg1)

## OLS estimation, Dep. Var.: fulltime_num
## Observations: 798
## Standard-errors: Heteroskedasticity-robust
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 9.99342   1.21845 8.20172 9.5089e-16 ***
## state       -2.26949   1.29598 -1.75118 8.0301e-02 .
## post        -2.27342   1.56537 -1.45232 1.4681e-01
## state:post   2.99498   1.68446  1.77801 7.5786e-02 .
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 8.24468 Adj. R2: 0.002734
```

Card and Krueger (1994) - Minimum wage and employment

```
reg1 <- feols(fulltime_num ~ state + post + state:post, data = ckdata, vcov = "HC1")
reg2 <- feols(parttime_num ~ state + post + state:post, data = ckdata, vcov = "HC1")
table <- etable(reg1, reg2,
                 # standard errors, digits, fit statistics, put SE below coefficients (the norm)
                 vcov = "HC1", digits = 3, fitstat = "", se.below = TRUE,
                 # change significance codes to the norm
                 signif.code = c("***" = 0.01, "**" = 0.05, "*" = 0.1),
                 # rename the variables
                 dict = c("Constant" = "Intercept", "state" = "Treat", "post" = "Post", "state:post" = "Treat x Post"))
table

##                                     reg1          reg2
## Dependent Var.: fulltime_num parttime_num
##
## Constant           9.99***      19.7*** 
##                   (1.22)        (1.09)  
## Treat            -2.27*       -1.04    
##                   (1.30)        (1.23)  
## Post             -2.27       0.103   
##                   (1.57)        (1.60)  
## Treat x Post     2.99*       -0.405  
##                   (1.68)        (1.80)  
## 
## -----
## S.E. type   Hetero.-rob. Hetero.-rob.
## --- 
## Signif. codes: 0 '***' 0.01 '**' 0.05 '*' 0.1 ' ' 1
# drop some rows
table <- table[-c(1:2, 11:nrow(table)),]
# rename columns
colnames(table) <- c("", "Full-time", "Part-time")
```

Card and Krueger (1994) - Minimum wage and employment

```
kable(table,
  align = "lcc", booktabs = TRUE, linesep = "", escape = FALSE, row.names = FALSE) %>%
  column_spec(1, width = "2cm") %>%
  column_spec(c(2:3), width = "1.5cm") %>%
  kable_styling() %>%
  footnote("* p < 0.1, ** p < 0.05, *** p < 0.01.", general_title = "",
    footnote_as_chunk = TRUE,
    escape = FALSE
  ) %>%
  footnote("Note: Robust standard errors in parentheses.", general_title = "",
    footnote_as_chunk = TRUE,
    escape = FALSE
  )
```

	Full-time	Part-time
Constant	9.99*** (1.22)	19.7*** (1.09)
Treat	-2.27* (1.30)	-1.04 (1.23)
Post	-2.27 (1.57)	0.103 (1.60)
Treat x Post	2.99* (1.68)	-0.405 (1.80)

Note: Robust standard errors in parentheses.

* p < 0.1, ** p < 0.05, *** p < 0.01.

Card and Krueger (1994) - Poisson regression!

```
reg1 <- feglm(fulltime_num ~ state + post + state:post, data = ckdata, vcov = "HC1", family = "poisson")
reg2 <- feglm(parttime_num ~ state + post + state:post, data = ckdata, vcov = "HC1", family = "poisson")
table <- etable(reg1, reg2,
                 # standard errors, digits, fit statistics, put SE below coefficients (the norm)
                 vcov = "HC1", digits = 3, fitstat = "", se.below = TRUE,
                 # change significance codes to the norm
                 signif.code = c("***" = 0.01, **" = 0.05, *" = 0.1),
                 # rename the variables
                 dict = c("Constant" = "Intercept", "state" = "Treat", "post" = "Post", "state:post" = "Treat x Post"))
table

##                                reg1          reg2
## Dependent Var.: fulltime_num parttime_num
##
## Constant           2.30***       2.98*** 
##                   (0.122)      (0.056)
## Treat            -0.258*       -0.054  
##                   (0.135)      (0.063)
## Post             -0.258        0.005  
##                   (0.176)      (0.081)
## Treat x Post     0.347*       -0.021  
##                   (0.192)      (0.092)
## 
## -----
## S.E. type    Hetero.-rob. Hetero.-rob.
## --- 
## Signif. codes: 0 '***' 0.01 '**' 0.05 '*' 0.1 ' ' 1
# drop some rows
table <- table[-c(1:2, 11:nrow(table)),]
# rename columns
colnames(table) <- c("", "Full-time", "Part-time")
```

Card and Krueger (1994) - Minimum wage and employment

```
kable(table, caption = "Poisson regression", # adding a caption
      align = "lcc", booktabs = TRUE, linesep = "", escape = FALSE, row.names = FALSE) %>%
  column_spec(1, width = "2cm") %>%
  column_spec(c(2:3), width = "1.5cm") %>%
  kable_styling() %>%
  footnote("* p < 0.1, ** p < 0.05, *** p < 0.01.", general_title = "",
          footnote_as_chunk = TRUE,
          escape = FALSE
         ) %>%
  footnote("Note: Robust standard errors in parentheses.", general_title = "",
          footnote_as_chunk = TRUE,
          escape = FALSE
         )
```

Table 1: Poisson regression

	Full-time	Part-time
Constant	2.30*** (0.122)	2.98*** (0.056)
Treat	-0.258* (0.135)	-0.054 (0.063)
Post	-0.258 (0.176)	0.005 (0.081)
Treat x Post	0.347* (0.192)	-0.021 (0.092)

Note: Robust standard errors in parentheses.

* p < 0.1, ** p < 0.05, *** p < 0.01.

Estimating diff-in-diff empirically - adding controls

- ▶ Can add control variables

$$Y_{it} = \beta_0 + \beta_1 TREAT_i + \beta_2 POST_t + \quad (2)$$

$$\beta_3(POST_t \times TREAT_i) + X_{it} + \varepsilon_{it} \quad (3)$$

- ▶ Adding controls can help control for differing trends (“conditional” parallel trends)
- ▶ Note: the interpretation of β_0 is no longer the same; others stay the same

Standard errors in differences-in-differences

- ▶ Card and Krueger did not cluster standard errors
 - ▶ In fact, that would have been difficult because they really only had two “clusters”!
 - ▶ But their robust standard errors are likely underestimated due to the clustering
- ▶ Classic reference: Bertrand, Duflo, and Mullainathan (2004)
 - ▶ “How Much Should We Trust Differences-in-Differences Estimates?”

Standard errors in differences-in-differences

- ▶ Bertrand, Duflo, and Mullainathan (2004) suggest three possibilities:
 - ① Cluster at the group level
 - ② Block bootstrap (not going to discuss)
 - ③ Aggregating data into one pre and one post period (event studies later)
- ▶ Let's go through these

- ▶ The most common approach: cluster standard errors
- ▶ Cameron, Gelbach, and Miller (2008) show that this is problematic with few clusters
 - ▶ “Bootstrap-based improvements for inference with clustered errors”
- ▶ The authors look at many possible approaches and find that the “wild cluster bootstrap” seems to perform best, on average

Wild cluster bootstrap

- ▶ The wild cluster bootstrap is a “non-parametric” bootstrap
 - ▶ I'll do the non-cluster as an example. Software makes this easy!
- ▶ Suppose we are interested in the following regression:

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$$

- ▶ We want to test whether $\beta_1 = 0$ and we have relatively few clusters (say between 5 and 30)

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i \quad (4)$$

- ① Estimate above regression and obtain $\hat{\beta}, \hat{\varepsilon}$
- ② Impose the null hypothesis ($\beta_1 = 0$) and estimate the restricted regression:

$$\tilde{y}_i = \tilde{\beta}_0 + \tilde{\varepsilon}_i \quad (5)$$

Wild cluster bootstrap

- ③ Bootstrap replications:
 - ▶ Use equation 5 to generate \tilde{y}_i^b , where $\tilde{y}_i^b = \tilde{\beta}_0^b + \tilde{\varepsilon}_i^b$
 - ▶ Rademacher weights: The randomness comes from either adding $\hat{\varepsilon}_i$ or $-\hat{\varepsilon}_i$ with equal probability
 - ▶ Estimate:

$$\tilde{y}_i^b = \tilde{\beta}_0^{b*} + \tilde{\beta}_1^{b*}x_i + \tilde{\varepsilon}_i^{b*} \quad (6)$$

- ▶ Calculate the t-statistic for the bootstrap replication:

$$t^{b*} = \frac{\tilde{\beta}_1^{b*}}{\sqrt{\tilde{V}^{b*}}} \quad (7)$$

Two-tailed test:

- ▶ Reject the null hypothesis if

$$|t^{b*}| > |t^*| \text{ for } b = 1, \dots, B,$$

where t^* is the t-statistic from the *original regression*.

- ▶ P-value across B bootstrap samples is:

$$\frac{1}{B} \sum_{b=1}^B \mathbb{I}(|t^{b*}| > |t^*|), \quad (8)$$

where \mathbb{I} is the indicator function.

Implementing WCB in R

- ▶ Thankfully there's a package that allows us to do this!
 - ▶ `fwildclusterboot` (Friedrich, 2019)
- ▶ This package works with `fixest` objects!
- ▶ Let's use the `castle.dta` data in the GitHub repo to test this

Implementing WCB in R

```
library(haven) # to load .dta files
df <- read_dta("castle.dta")
head(df)

## # A tibble: 6 x 185
##   state    year    sid    cdl pre2_cdl caselaw anywhere assumption civil homicide_c
##   <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Alaba~ 2000     1     0     0     0     0     0     0     0     329
## 2 Alaba~ 2001     1     0     0     0     0     0     0     0     379
## 3 Alaba~ 2002     1     0     0     0     0     0     0     0     303
## 4 Alaba~ 2003     1     0     0     0     0     0     0     0     299
## 5 Alaba~ 2004     1     0     1     0     0     0     0     0     254
## 6 Alaba~ 2005     1     0     1     0     0     0     0     0     374
## # i 175 more variables: robbery_gun_r <dbl>, jhcitizen_c <dbl>,
## #   jhpolice_c <dbl>, homicide <dbl>, robbery <dbl>, assault <dbl>,
## #   burglary <dbl>, larceny <dbl>, motor <dbl>, murder <dbl>,
## #   hc_felonywsus <dbl>, jhcitizen <dbl>, jhpolice <dbl>, population <dbl>,
## #   police <dbl>, unemployrt <dbl>, income <dbl>, blackm_15_24 <dbl>,
## #   whitem_15_24 <dbl>, blackm_25_44 <dbl>, whitem_25_44 <dbl>, prisoner <dbl>,
## #   lagprisoner <dbl>, poverty <dbl>, exp_subsidy <dbl>, ...
# key variables: state, year, cdl ("treatment"), and homicide_c (outcome)
# homicide_c to rate (per 100,000 people)
df$homicide_c <- (df$homicide_c/df$population)*100000
# and log
df$homicide_c <- log(df$homicide_c)
```

Implementing WCB in R

```
# Note: this is not differences-in-differences.  
# Just an example of the wild cluster bootstrap with fwildclusterboot  
reg1 <- feols(homicide_c ~ cdl, data = df, cluster = "state")  
summary(reg1)  
  
## OLS estimation, Dep. Var.: homicide_c  
## Observations: 550  
## Standard-errors: Clustered (state)  
##             Estimate Std. Error t value Pr(>|t|)  
## (Intercept) 1.360284  0.082267 16.53498 < 2.2e-16 ***  
## cdl        0.338002  0.092532  3.65279 0.00063065 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
## RMSE: 0.578226  Adj. R2: 0.036511
```

Implementing WCB in R

```
reg1 <- feols(homicide_c ~ cdl, data = df, cluster = "state")
boot_reg <- boottest(
    reg1,
    clustid = c("state"),
    param = "cdl",
    B = 10000,
    type = "rademacher" # default weighting, by the way
)
boot_reg

## boottest.fixest(object = reg1, param = "cdl", B = 10000, clustid = c("state"),
##   type = "rademacher")
##
## p value: 2e-04
## confidence interval: 0.1536 0.522
## test statistic 3.6528
```

Add controls

```
reg1 <- feols(homicide_c ~ cdl + log(population) + log(income) + unemployrt, data = df, cluster = "state")
reg1

## OLS estimation, Dep. Var.: homicide_c
## Observations: 550
## Standard-errors: Clustered (state)
##             Estimate Std. Error   t value Pr(>|t|)
## (Intercept) 8.328111  5.001764 1.665035 1.0229e-01
## cdl         0.156808  0.088948 1.762923 8.4149e-02 .
## log(population) 0.304923  0.062867 4.850276 1.2902e-05 ***
## log(income)    -1.061536  0.477547 -2.222894 3.0868e-02 *
## unemployrt    -0.006588  0.021765 -0.302699 7.6340e-01
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.471373  Adj. R2: 0.35618

boot_reg <- boottest(
  reg1,
  clustid = c("state"),
  param = "cdl",
  B = 10000,
  type = "rademacher" # default weighting, by the way
)
boot_reg

## boottest.fixest(object = reg1, param = "cdl", B = 10000, clustid = c("state"),
##                 type = "rademacher")
## 
## p value: 0.0979
## confidence interval: -0.0308 0.3425
## test statistic 1.7629
```

Multi-way clustering, too!

```
reg1 <- feols(homicide_c ~ cdl + log(population) + log(income) + unemployrt, data = df, cluster = c("state", "year"))
reg1

## OLS estimation, Dep. Var.: homicide_c
## Observations: 550
## Standard-errors: Clustered (state & year)
##             Estimate Std. Error   t value Pr(>|t|)
## (Intercept)  8.328111  5.024169  1.657609 0.12838986
## cdl         0.156808  0.100698  1.557209 0.15047711
## log(population) 0.304923  0.062555  4.874462 0.00064729 ***
## log(income)    -1.061536  0.476525 -2.227661 0.05004055 .
## unemployrt    -0.006588  0.024931 -0.264260 0.79694594
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.471373  Adj. R2: 0.35618

boot_reg <- boottest(
  reg1,
  clustid = c("state"),
  param = "cdl",
  B = 10000,
  type = "rademacher" # default weighting, by the way
)
boot_reg

## boottest.fixest(object = reg1, param = "cdl", B = 10000, clustid = c("state"),
##                 type = "rademacher")
## 
## p value: 0.098
## confidence interval: -0.0305 0.3397
## test statistic 1.7629
```

Finally, Webb weights (Webb, 2023) – but using Rademacher weights is the norm

```
reg1 <- feols(homicide_c ~ cdl + log(population) + log(income) + unemployrt, data = df, cluster = c("state", "year"))
reg1

## OLS estimation, Dep. Var.: homicide_c
## Observations: 550
## Standard-errors: Clustered (state & year)
##             Estimate Std. Error   t value Pr(>|t|)
## (Intercept) 8.328111  5.024169  1.657609 0.12838986
## cdl         0.156808  0.100698  1.557209 0.15047711
## log(population) 0.304923  0.062555  4.874462 0.00064729 ***
## log(income)    -1.061536  0.476525 -2.227661 0.05004055 .
## unemployrt    -0.006588  0.024931 -0.264260 0.79694594
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.471373  Adj. R2: 0.35618

boot_reg <- boottest(
  reg1,
  clustid = c("state"),
  param = "cdl",
  B = 10000,
  type = "webb"
)
boot_reg

## boottest.fixest(object = reg1, param = "cdl", B = 10000, clustid = c("state"),
##                 type = "webb")
## 
## p value: 0.0967
## confidence interval: -0.0309 0.3376
## test statistic 1.7629
```

Some thoughts on clustering

- ▶ If you have more than ~30 clusters, you can probably just cluster at the group level
 - ▶ We see that the standard errors are very similar in the state examples above
- ▶ Otherwise, consider using an alternative approach
 - ▶ Also important when clusters have wildly different sample sizes, or where the treated clusters are relatively few (Moulton, 1990)
- ▶ Alternative approach with only one treated cluster: randomization inference

- ▶ They are interested in the change in insurance coverage in Hawaii relative to other states:

$$Y_{ist} = X_{ist}\beta^t + Z_{st}\gamma^t + H_{it}\delta^t + \phi_{st} + \eta_{it} \quad (9)$$

- ▶ They calculate the change as: $\Delta = \delta^1 - \delta^0$
- ▶ The idea: see where the Hawaii effect sits in the distribution of the same effect across *all US states*
 - ▶ “Placebo” effects
 - ▶ Note that this is not true “randomization” inference
 - ▶ I’ll show you an example with one of my papers in a minute

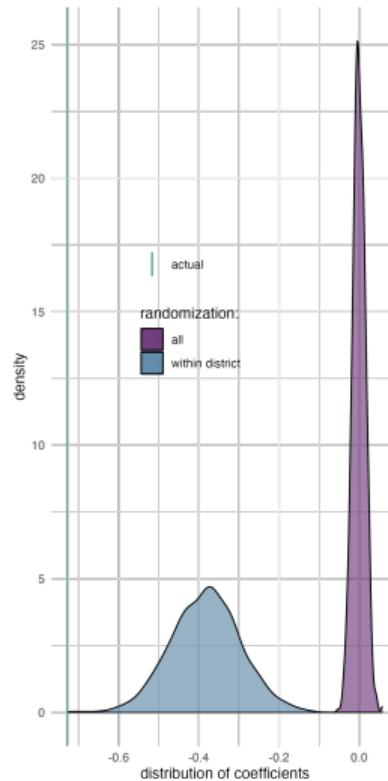
Randomization inference in Buchmueller, DiNardo, and Valletta (2011)

TABLE 1—STATE EFFECTS (*percentages*), OWN ESI COVERAGE (by quintile)
(Placebo tests, Hawaii versus all other states)

	Full sample	ESI quintiles				
		1st	2nd	3rd	4th	5th
<i>Panel A. 1979–1982</i>						
Hawaii effect	0.139*	0.209*	0.133*	0.132**	0.085**	0.002
Placebo effects (other states)						
95th percentile	0.046	0.066	0.065	0.055	0.045	0.031
5th percentile	−0.046	−0.062	−0.070	−0.070	−0.052	−0.046
<i>Panel B. 2002–2005</i>						
Hawaii effect	0.181**	0.281**	0.232**	0.142**	0.101**	0.031
Placebo tests (other states)						
95th percentile	0.043	0.045	0.062	0.054	0.045	0.039
5th percentile	−0.057	−0.067	−0.069	−0.070	−0.054	−0.052
<i>Panel C. Difference (2002–2005 minus 1979–1982)</i>						
Hawaii effect	0.042	0.071*	0.099*	0.011	0.016	0.030
Placebo tests (other states)						
95th percentile	0.051	0.067	0.087	0.077	0.052	0.049
5th percentile	−0.046	−0.076	−0.065	−0.066	−0.068	−0.040

- ▶ I'm interested in the effects of pollution on agricultural productivity in India
- ▶ I have villages, which are nested within districts
 - ▶ I cluster on villages
- ▶ Alternative: randomly assign pollution to villages *within the same district* and compare my effects to the distribution of effects

Randomization inference in Merfeld (2023)



- ▶ Above, we looked at the parallel trends assumption graphically in Richardson and Troost (2009)
- ▶ Another common way is to look at *leads* of treatment
 - ▶ In my paper, for example, pollution next year should not affect agricultural productivity this year

Leads of pollution in Merfeld (2023)

	(1)	(2)
particulate matter (one-year lead)	-0.033 (0.067)	
particulate matter (two-year lead)		-0.070 (0.052)
weather (expanded)	No	No
fixed effects:		
village	Yes	Yes
year	Yes	Yes
F	592	783
observations	1,161,265	1,055,562

Note: Standard errors are in parentheses and are clustered at the village level. *

p<0.10 ** p<0.05 *** p<0.01

Convincing the reader is like writing a good story

- ▶ When you're writing a diff-in-diff paper, think about the possible threats to your identification strategy
- ▶ Then, think about how you can convince the reader that your strategy is valid
 - ▶ Use placebos: is there somewhere we shouldn't expect an effect?
 - ▶ In the case of my paper, the leads convinced some seminar participants!
- ▶ You can likewise think of heterogeneity we would *expect* to see, and test for that!

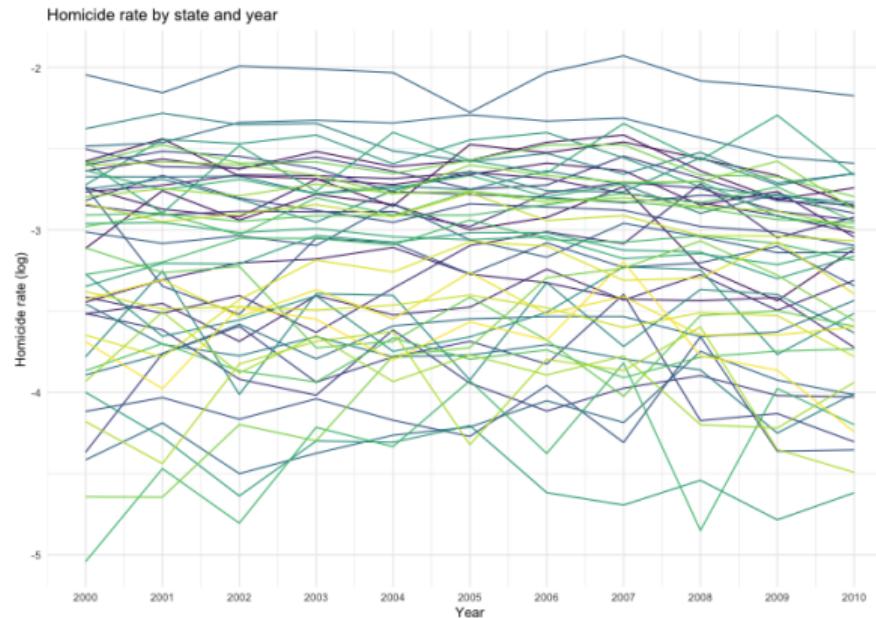
Before moving on to some of the new literature...

- ▶ Let's talk about fixed and random effects
 - ▶ Fixed effects will be important for the upcoming discussions
- ▶ Some nice (but old) slides from Oscar Tores-Reyna [here](#).

Panel data

- ▶ Both fixed and random effects are used in panel data
 - ▶ Panel data: data with multiple observations for each unit
 - ▶ Examples: individuals, firms, countries, etc.
- ▶ In our previous example of homicide and castle doctrine laws: unit is the state!

Panel data



Fixed effects

- ▶ Fixed effects are a way to control for omitted variables
 - ▶ However there is a key assumption: the omitted variable is time-invariant
- ▶ Fixed effects are also called “within” effects
 - ▶ Why? Because we are looking at the variation within each unit
- ▶ The regression is of the form:

$$y_{it} = \alpha_i + \beta x_{it} + \varepsilon_{it}, \quad (10)$$

where α_i is the fixed effect for unit i . Note the subscript! No t!

- ▶ Empirically, what are fixed effects doing?
 - ▶ They are subtracting the mean of **each unit** from the outcome variable
- ▶ In a regression, we add a dummy variable for each unit
 - ▶ We have to leave out one dummy variable, though
 - ▶ Software will do this for us!
 - ▶ Note that the intercept is usually meaningless in this case
- ▶ Cannot include time-invariant variables in the regression
 - ▶ Why? Because the fixed effect will absorb them!

Fixed effects with feols

► feols makes this easy on us. Let's return to our previous example.

```
reg1 <- feols(homicide_c ~ cdl + log(population) + log(income) + unemployrt, data = df, cluster = c("state"))
reg2 <- feols(homicide_c ~ cdl + log(population) + log(income) + unemployrt | state, data = df, cluster = c("state"))
etable(reg1, reg2,
      # standard errors, digits, fit statistics, put SE below coefficients (the norm)
      digits = 3, fitstat = "", se.below = TRUE,
      # change significance codes to the norm
      signif.code = c("***" = 0.01, "**" = 0.05, "*" = 0.1))

##                      reg1      reg2
## Dependent Var.: homicide_c homicide_c
##
## Constant          3.72
##                  (5.00)
## cdl             0.157*
##                  (0.089)   0.059
## log(population) 0.305*** -0.151
##                  (0.063)   (0.341)
## log(income)     -1.06**
##                  (0.478)   0.242
## unemployrt      -0.007   -0.024***
##                  (0.022)   (0.006)
## Fixed-Effects: -----
## state            No       Yes
## -----
## S.E.: Clustered by: state by: state
## ---
## Signif. codes: 0 *** 0.01 ** 0.05 * 0.1 ' ' 1
```

Just a quick note: wild cluster bootstrap still works!

```
# need to use a numeric variable for the bootstrap. sid is in our data, thankfully.
reg1 <- feols(homicide_c ~ cdl + log(population) + log(income) + unemployrt | sid, data = df, cluster = c("sid"))
reg1

## OLS estimation, Dep. Var.: homicide_c
## Observations: 550
## Fixed-effects: sid: 50
## Standard-errors: Clustered (sid)
##             Estimate Std. Error   t value Pr(>|t|)
## cdl          0.058781  0.049370  1.190624 0.23953723
## log(population) -0.151406  0.341426 -0.443451 0.65939131
## log(income)      0.241548  0.335153  0.720711 0.47451204
## unemployrt     -0.023468  0.006455 -3.635678 0.00066447 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.179187   Adj. R2: 0.897774
##             Within R2: 0.058749
boot_reg <- boottest(
  reg1,
  clustid = c("sid"), # note that it requires a numeric variable!
  param = "cdl",
  B = 10000
)
boot_reg

## boottest.fixest(object = reg1, param = "cdl", B = 10000, clustid = c("sid"))
##
## p value: 0.277
## confidence interval: -0.033 0.1601
## test statistic 1.1917
```

- ▶ It's not quite the same as the canonical differences-in-differences model
- ▶ We redefine treatment for the same unit
 - ▶ Before treatment, the value is zero, and after it is One
 - ▶ Note that this is different from the canonical model, where the value is zero for the comparison group and one for the treated group
- ▶ In fact, the regression we just saw is a differences-in-differences model of this form!
 - ▶ In practice, we often tend to add time fixed effects, too:

$$y_{it} = \alpha_i + \delta_t + \beta x_{it} + \varepsilon_{it}, \quad (11)$$

The “effect” of castle doctrine laws, two-way fixed effects

```
reg1 <- feols(homicide_c ~ cdl + log(population) + log(income) + unemployrt | state, data = df, cluster = c("state"))
reg2 <- feols(homicide_c ~ cdl + log(population) + log(income) + unemployrt | state + year, data = df, cluster = c("state"))
etable(reg1, reg2,
      # standard errors, digits, fit statistics, put SE below coefficients (the norm)
      digits = 3, fitstat = "", se.below = TRUE,
      # change significance codes to the norm
      signif.code = c("***" = 0.01, "**" = 0.05, "*" = 0.1))

##                               reg1      reg2
## Dependent Var.: homicide_c homicide_c
##
##   cdl          0.059      0.076
##             (0.049)    (0.056)
##   log(population) -0.151     -0.974**
##             (0.341)    (0.484)
##   log(income)    0.242      0.283
##             (0.335)    (0.308)
##   unemployrt   -0.024***   -0.009
##             (0.006)    (0.012)
## Fixed-Effects: -----
##   state        Yes       Yes
##   year         No       Yes
## -----
## S.E.: Clustered by: state by: state
## ---
## Signif. codes: 0 *** 0.01 ** 0.05 * 0.1 ' ' 1
```

- ▶ Before turning to recent issues discovered with the two-way fixed effects estimator, let's talk about random effects
- ▶ Random effects are a way to capture the heterogeneity across units
 - ▶ The key is that this heterogeneity is *random* and uncorrelated with the predictors in the model
- ▶ This is really a way to capture the *variance* across units
 - ▶ In practice, this absorbs some of the variance, increasing precision (but at the cost of the assumption above)

Random effects in R

```
library(lme4)
reg1 <- feols(homicide_c ~ cdl + log(population) + log(income) + unemployrt, data = df, cluster = "state")
# note iid standard errors for simplicity for random effects
# (can use other packages to change the vcov calculation if we think assumptions aren't exactly true...)
reg2 <- lmer(homicide_c ~ cdl + log(population) + log(income) + unemployrt + (1 | state), data = df)
reg3 <- feols(homicide_c ~ cdl + log(population) + log(income) + unemployrt | state, data = df, cluster = "state")
modelsummary(list("None" = reg1, "RE" = reg2, "FE" = reg3), gof_omit = ".*",
            output = "markdown", coef_omit = c(-2,-3,-4,-5),
            coef_rename = c("cdl" = "Castle laws", "log(population)" = "Population (log)",
                           "log(income)" = "Income (log)", "unemployrt" = "Unemployment rate"))
```

	None	RE	FE
Castle laws	0.157 (0.089)	0.043 (0.029)	0.059 (0.049)
Population (log)	0.305 (0.063)	0.276 (0.066)	-0.151 (0.341)
Income (log)	-1.062 (0.478)	0.001 (0.191)	0.242 (0.335)
Unemployment rate	-0.007 (0.022)	-0.029 (0.006)	-0.023 (0.006)

Some things about random effects relative to vanilla OLS (not fixed effects):

- ▶ The random effects estimator is asymptotically more efficient than OLS if there is unit-level heterogeneity: $\mathbb{E}(V_{RE}) < \mathbb{E}(V_{OLS})$
 - ▶ In practice with finite samples, not necessarily
- ▶ In expectation, coefficients are the same: $\mathbb{E}(\beta_{RE}) = \mathbb{E}(\beta_{OLS})$
 - ▶ Random effects is estimated using (feasible) generalized least squares, which essentially reweights the observations

- ▶ Recently, a number of papers have shown that the two-way fixed effects estimator can be... problematic
- ▶ We have been discussing differences-in-differences with TWFE of the following form:

$$y_{it} = \alpha_i + \delta_t + \beta D_{it} + \gamma x_{it} + \varepsilon_{it}, \quad (12)$$

where D_{it} is a dummy variable for treatment.

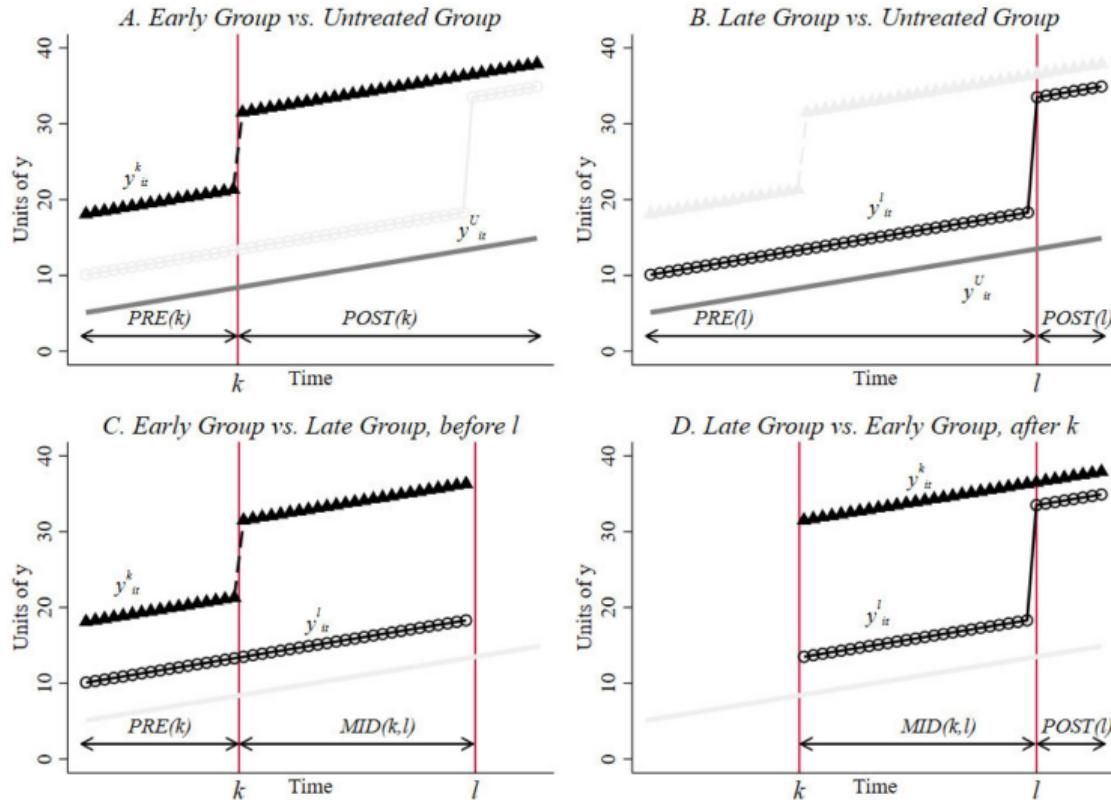
- ▶ If there are only two time periods and one group receives treatment in only one period, this is not a problem!
 - ▶ The Card and Krueger setup is not an issue with TWFE

- ▶ The real issue is when treatment is staggered across time
 - ▶ For example, if treatment is introduced in different years for different states
- ▶ It turns out this is the case with the castle doctrine law!

year	treated
2000	0.00
2001	0.00
2002	0.00
2003	0.00
2004	0.00
2005	0.00
2006	0.02
2007	0.28
2008	0.36
2009	0.40
2010	0.42

- ▶ Goodman-Bacon lays out the problem (as does Scott in *CI*)
- ▶ Suppose we have three groups and three time periods
 - ▶ Group 1 is treated before period 2 (Goodman-Bacon calls this group k)
 - ▶ Group 2 is treated before period 3 (Goodman-Bacon calls this group l)
 - ▶ Group 3 is never treated (Goodman-Bacon calls this group U)
- ▶ If we are willing to assume treatment effect *homogeneity*, then we have no problems!
 - ▶ Are you willing to assume this?

- ▶ Goodman-Bacon shows that the overall treatment effect is a *weighted average* of treatment effects from every possible 2x2 comparison where treatment status doesn't change:
 - ▶ Group 1 vs group 2
 - ▶ Group 1 vs group 3
 - ▶ Group 2 vs group 1
 - ▶ Group 2 vs group 3
- ▶ These weights are a function of two things:
 - ▶ Group sizes
 - ▶ Variance in treatment



- He shows there are three relevant comparisons:

$$\hat{\beta}_{kU}^{2\times 2} \equiv (\bar{y}_k^{POST(k)} - \bar{y}_k^{PRE(k)}) - (\bar{y}_U^{POST(k)} - \bar{y}_U^{PRE(k)}) \quad (13)$$

$$\hat{\beta}_{kl}^{2\times 2,k} \equiv (\bar{y}_k^{MID(k,l)} - \bar{y}_k^{PRE(k)}) - (\bar{y}_l^{MID(k,l)} - \bar{y}_l^{PRE(k)}) \quad (14)$$

$$\hat{\beta}_{kl}^{2\times 2,l} \equiv (\bar{y}_l^{POST(l)} - \bar{y}_l^{MID(k,l)}) - (\bar{y}_k^{POST(l)} - \bar{y}_k^{MID(k,l)}), \quad (15)$$

where k and l are treated groups, and U is the untreated group.

- ▶ The DD estimator is a weighted average of all these comparisons.
- ▶ Generalizing to K time periods:

$$\hat{\beta}^{DD} = \sum_{k \neq U} s_{kU} \hat{\beta}_{kU}^{2 \times 2} + \sum_{k \neq U} \sum_{l > k} \left[s_{kl}^k \hat{\beta}_{kl}^{2 \times 2, k} + s_{kl}^l \hat{\beta}_{kl}^{2 \times 2, l} \right], \quad (16)$$

where s_{ij} is the weight for the comparison between groups i and j .

The weights:

$$s_{kU} = \frac{(n_k + n_U)^2 n_{kU} (1 - n_{kU}) \bar{D}_k (1 - \bar{D}_k)}{\hat{V}^D} \quad (17)$$

$$s_{kl}^k = \frac{\left((n_k + n_l) (1 - \bar{D}_l) \right)^2 n_{kl} (1 - n_{kl}) \frac{\bar{D}_k - \bar{D}_l}{1 - \bar{D}_l} \frac{1 - \bar{D}_k}{1 - \bar{D}_k}}{\hat{V}^D} \quad (18)$$

$$s_{kl}^l = \frac{\left((n_k + n_l) (\bar{D}_k) \right)^2 n_{kl} (1 - n_{kl}) \frac{\bar{D}_l}{\bar{D}_k} \frac{\bar{D}_k - \bar{D}_l}{\bar{D}_k}}{\hat{V}^D} \quad (19)$$

- ▶ Note how the variance of treatment affects the weights! “Changing the number or spacing of time periods changes the weights” (Goodman-Bacon).
 - ▶ Even if the treatment effect is constant, changing the length of the panel can change the weighted average if different groups have different treatment effects.

- ▶ de Chaisemartin and D'Haultfœuille (2020) more explicitly show the problem with weights.
- ▶ Let $w_{g,t}$ represent the weight for group g at time t , $\Delta_{g,t}$ represent the average treatment effect for the same, and N_1 represent the total number of treated units at the time:

$$\beta_{FE} = \mathbb{E} \left[\sum_{(g,t): D_{g,t}=1} \frac{N_{g,t}}{N_1} w_{g,t} \Delta_{g,t} \right] \quad (20)$$

- ▶ The fixed effect estimator is a weighted average of these treatment effects.
 - ▶ But how do we get the weights?

- ▶ Let $D_{g,t}$ represent treatment status of group g at time t :

$$D_{g,t} = \alpha + \gamma_g + \lambda_t + \varepsilon_{g,t}, \quad (21)$$

where γ_g is the group fixed effect and λ_t is the time fixed effect.

- ▶ It turns out that the average residual *is not zero*, so let's rescale it by its mean:

$$w_{g,t} = \frac{\varepsilon_{g,t}}{\sum_{(g,t):D_{g,t}=1} \frac{N_{g,t}}{N_1} \varepsilon_{g,t}} \quad (22)$$

- ▶ Since $\varepsilon_{g,t}$ can be positive or negative, this means that the weights can be negative!

- ▶ The problem is an *extrapolation* problem
 - ▶ Essentially, “the regression predicts a treatment probability larger than the one in that cell” (de Chaisemartin and D'Haultfœuille)
- ▶ Note that if the treatment effect is constant, then the weighted average is always the same, no matter the weights
 - ▶ But if the treatment effect is not constant, then the weighted average can be different from the true average

- ▶ Let's return to Bacon-Goodman's formulation
- ▶ What happens if we have a "control" group in a later period that is treated in an earlier period?

$$\begin{aligned}\hat{\delta}_{lk}^{2 \times 2} = & ATT_{I, POST(I)} \\ & + \Delta Y_I^0(Post(I), MID) - \Delta Y_k^0(Post(I), MID) \\ & - (ATT_k(Post) - ATT_k(Mid))\end{aligned}\tag{23}$$

- ▶ The first line is *what we want*
- ▶ The second line is parallel-trends bias
- ▶ The third line is bias due to heterogeneity in time!
 - ▶ Even with parallel trends, this third line can cause deviations from the true ATT

Enough math, what to do?

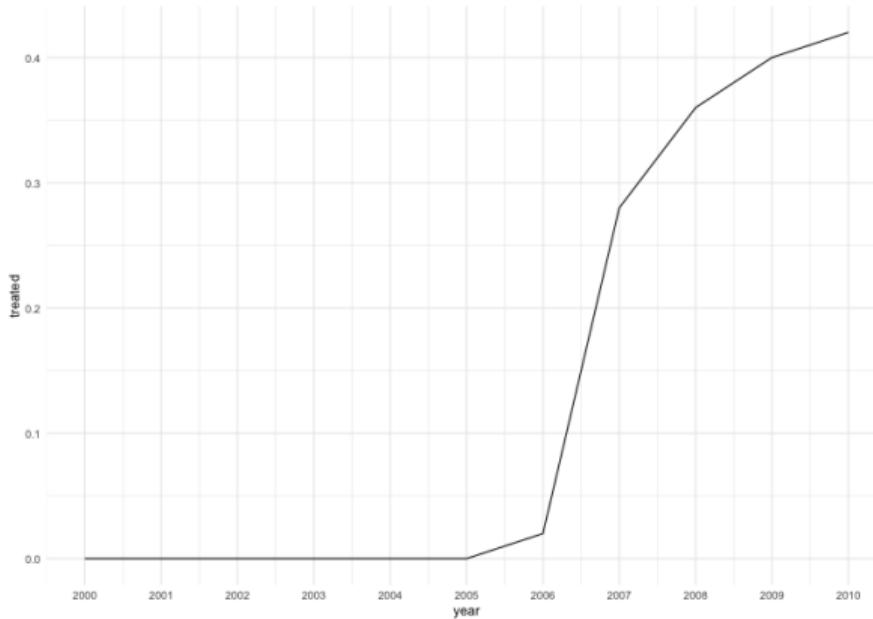
- ▶ That's enough math. Let's talk about what we can actually do!
- ▶ Let's go back to the castle.dta dataset
 - ▶ Cheng and Hoekstra (2013), *Journal of Human Resources*
- ▶ Let's use the information in Cunningham's book

- ▶ The “castle doctrine” laws essentially make lethal force “more” legal
- ▶ Recall the changes we made: turn homicide into a rate (per 100,000 people) and take the log:

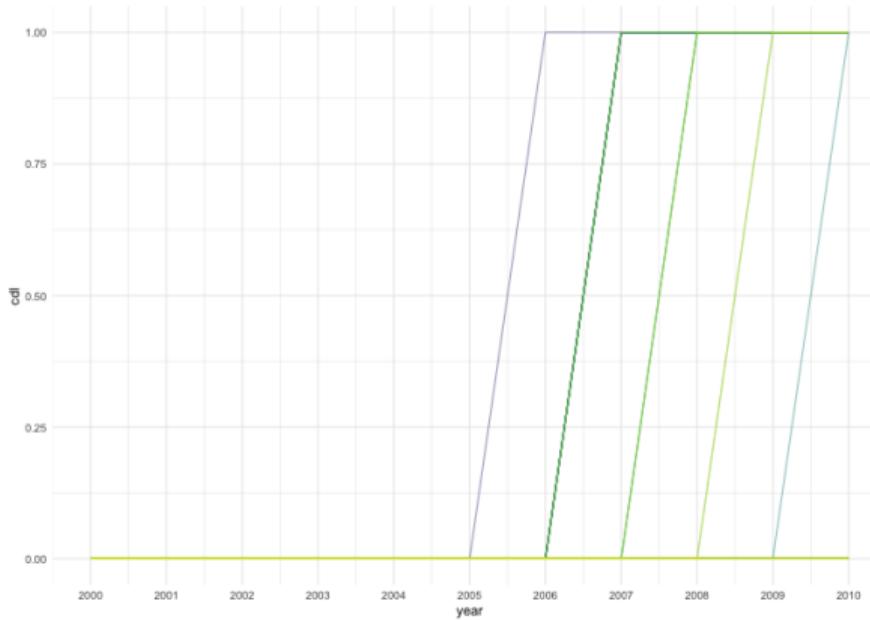
```
library(haven) # to load .dta files
df <- read_dta("castle.dta")
head(df)
# key variables: state, year, cdl ("treatment"), and homicide_c (outcome)
# homicide_c to rate (per 100,000 people)
df$homicide_c <- (df$homicide_c/df$population)*100000
# and log
df$homicide_c <- log(df$homicide_c)
```

- ▶ I made some changes to the data, as well
 - ▶ I've turned the "treatment" variable (cd1) into a dummy variable
- ▶ We have the issue we ran into above:
 - ▶ Treatment is staggered across time
 - ▶ This means that some already-treated units will serve as controls for later-treated units!

Treatment timing



Treatment timing by individual state



Two-way fixed effects with fixest, as simple as possible

```
# state fe, note the weights!
reg1 <- feols(homicide_c ~ cdl + log(population) + unemployrt | state, data = df,
               cluster = c("state"), weights = df$population)
# state and year fe
reg2 <- feols(homicide_c ~ cdl + log(population) + unemployrt | state + year, data = df,
               cluster = c("state"), weights = df$population)
# with state linear trends, note the syntax
reg3 <- feols(homicide_c ~ cdl + log(population) + unemployrt | state + year + state[year], data = df,
               cluster = c("state"), weights = df$population)
# put together
tablech <- etable(reg1, reg2, reg3,
                  # standard errors, digits, fit statistics, put SE below coefficients (the norm)
                  digits = 3, fitstat = "n", se.below = TRUE,
                  depvar = FALSE,
                  # rename the variables
                  dict = c("cdl" = "Treat", "log(population)" = "Pop. (log)", "unemployrt" = "Unemp. rate"))
tablech <- tablech[-c(12,13),]
tablech[c(7,10),2:4] <- ""
```

Two-way fixed effects with fixest

	reg1	reg2	reg3
Treat	0.057 (0.034)	0.089** (0.033)	0.097* (0.036)
Pop. (log)	-0.544** (0.168)	-0.768* (0.370)	-0.543 (0.960)
Unemp. rate	-0.024*** (0.004)	-0.009 (0.010)	-0.003 (0.011)
Fixed-Effects:			
state	Yes	Yes	Yes
year	No	Yes	Yes
Varying Slopes:			
year (state)	No	No	Yes
Observations	550	550	550

Note: Standard errors clustered at the state level are in parentheses.

* p < 0.1, ** p < 0.05, *** p < 0.01.

Event studies

- ▶ Event studies are a way to look at the effect of treatment over time
 - ▶ We can see if the effect is immediate, or if it takes time to “kick in”
 - ▶ We can also see whether there are any pre-trends
- ▶ Effectively, what we want to do is redefine the time period to be relative to treatment
 - ▶ For example, if treatment is introduced in 2005, we want to redefine 2005 as year 0, 2004 as year -1, etc.
 - ▶ Let's do this now

Event studies

```
# first, find the year of treatment by state
df <- df %>%
  # group by state ("panel" identifier)
  group_by(state) %>%
  mutate(year_treat = ifelse(cdl==1, year, NA),
    # first year with treatment
    year_treat = min(year_treat, na.rm = TRUE),
    # create new time variable called event_year
    event_year = year - year_treat) %>%
  # ungroup
  ungroup()
# note that states NEVER treated are missing
table(df$event_year)

## 
## -Inf   -10    -9    -8    -7    -6    -5    -4    -3    -2    -1     0     1     2     3     4
##  319      1      3      7     20     21     21     21     21     21     21     21     20     18     14     1

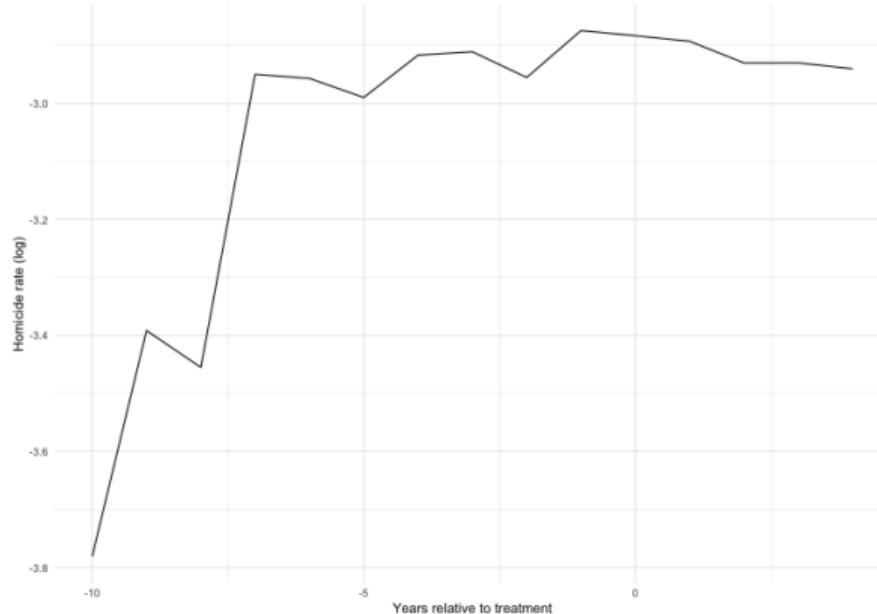
# now find the average homicide rate by event_year
df_event <- df %>%
  # drop the missings
  filter(event_year>-11) %>%
  group_by(event_year) %>%
  summarize(homicide_c = weighted.mean(homicide_c, weights = population, na.rm = TRUE))
```

check it looks okay

```
## # A tibble: 11 x 4                                ## # A tibble: 11 x 4
##   state    year event_year   cdl      state    year event_year   cdl
##   <chr>   <dbl>     <dbl> <dbl>    <chr>   <dbl>     <dbl> <dbl>
## 1 Alabama  2000      -7     0       1 Alaska   2000      -7     0
## 2 Alabama  2001      -6     0       2 Alaska   2001      -6     0
## 3 Alabama  2002      -5     0       3 Alaska   2002      -5     0
## 4 Alabama  2003      -4     0       4 Alaska   2003      -4     0
## 5 Alabama  2004      -3     0       5 Alaska   2004      -3     0
## 6 Alabama  2005      -2     0       6 Alaska   2005      -2     0
## 7 Alabama  2006      -1     0       7 Alaska   2006      -1     0
## 8 Alabama  2007       0     1       8 Alaska   2007       0     1
## 9 Alabama  2008       1     1       9 Alaska   2008       1     1
## 10 Alabama 2009       2     1      10 Alaska  2009       2     1
## 11 Alabama 2010       3     1      11 Alaska  2010       3     1
```

Plot the pure, *means*

```
ggplot(df_event) +  
  geom_line(aes(x = event_year, y = homicide_c)) +  
  theme_minimal() +  
  labs(x = "Years relative to treatment", y = "Homicide rate (log)")
```



What do we really want to see?

- ▶ We don't really want the means, though
- ▶ What do we want instead?

What do we really want to see?

- ▶ We don't really want the means, though
- ▶ What do we want instead?
 - ▶ We want the *effect* of treatment over time
 - ▶ We want to essentially plot *coefficients*

Calculating year-specific coefficients

```
# note that the year BEFORE treatment, -1, IS THE OMITTED CATEGORY
# i() is a fixedest-specific way to create factors/dummies
reg1 <- feols(homicide_c ~ i(event_year, ref = -1) + log(population) + unemployrt | state + year, data = df,
               cluster = c("state"), weights = df$population)
reg1$coefficients

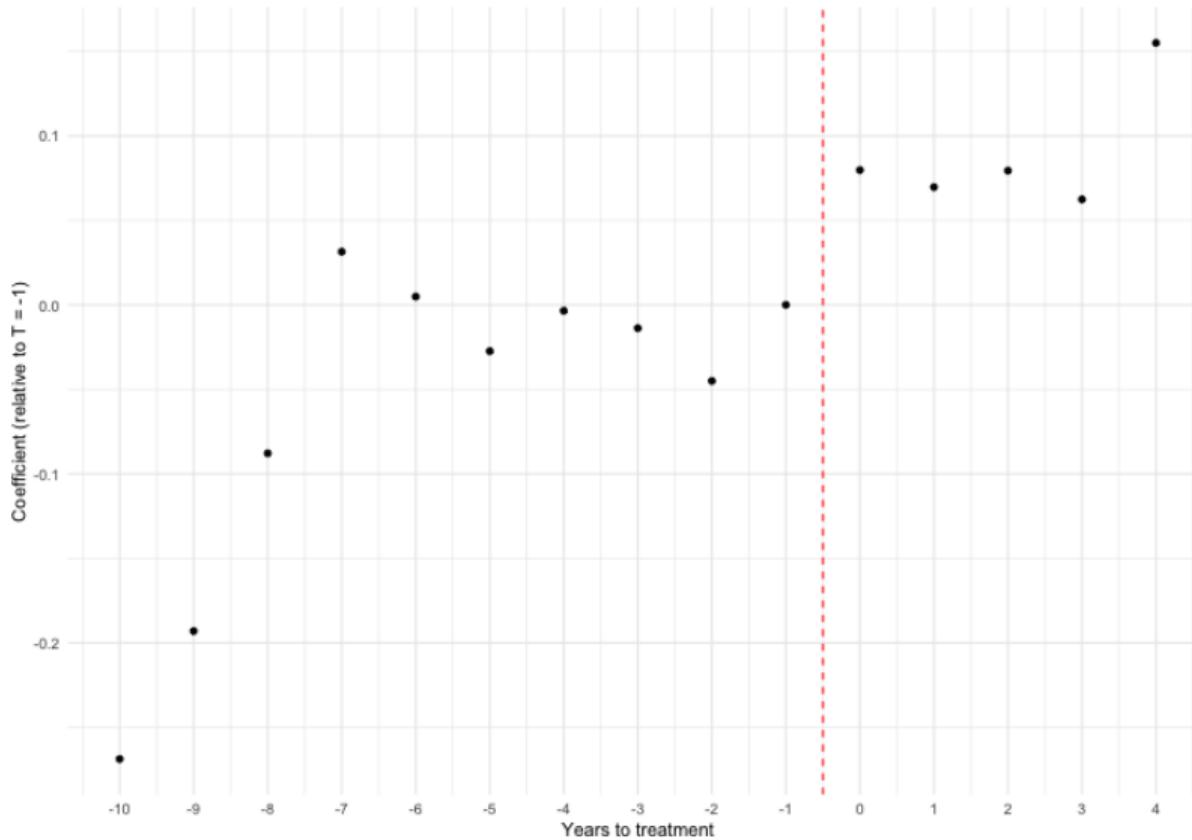
## event_year:::-10  event_year:::-9  event_year:::-8  event_year:::-7  event_year:::-6
##   -0.268714642    -0.193012613    -0.087832067     0.031459588     0.004916894
## event_year:::-5  event_year:::-4  event_year:::-3  event_year:::-2  event_year:::0
##   -0.027364799    -0.003498961    -0.013767586    -0.045036279     0.079746416
## event_year:::1  event_year:::2  event_year:::3  event_year:::4 log(population)
##   0.069676912     0.079426337     0.062440521     0.155021385    -0.724793998
## unemployrt
##   -0.009803301

# It's a vector. We can extract the coefficients we want by subsetting with []
# get coefficients
coef <- c(reg1$coefficients[1:9], 0, reg1$coefficients[10:14])
# confidence intervals
lower <- c(confint(reg1)[1:9,1], 0, confint(reg1)[10:14,1])
upper <- c(confint(reg1)[1:9,2], 0, confint(reg1)[10:14,2])
# create minimum/maximum
years <- c(-10:4)
```

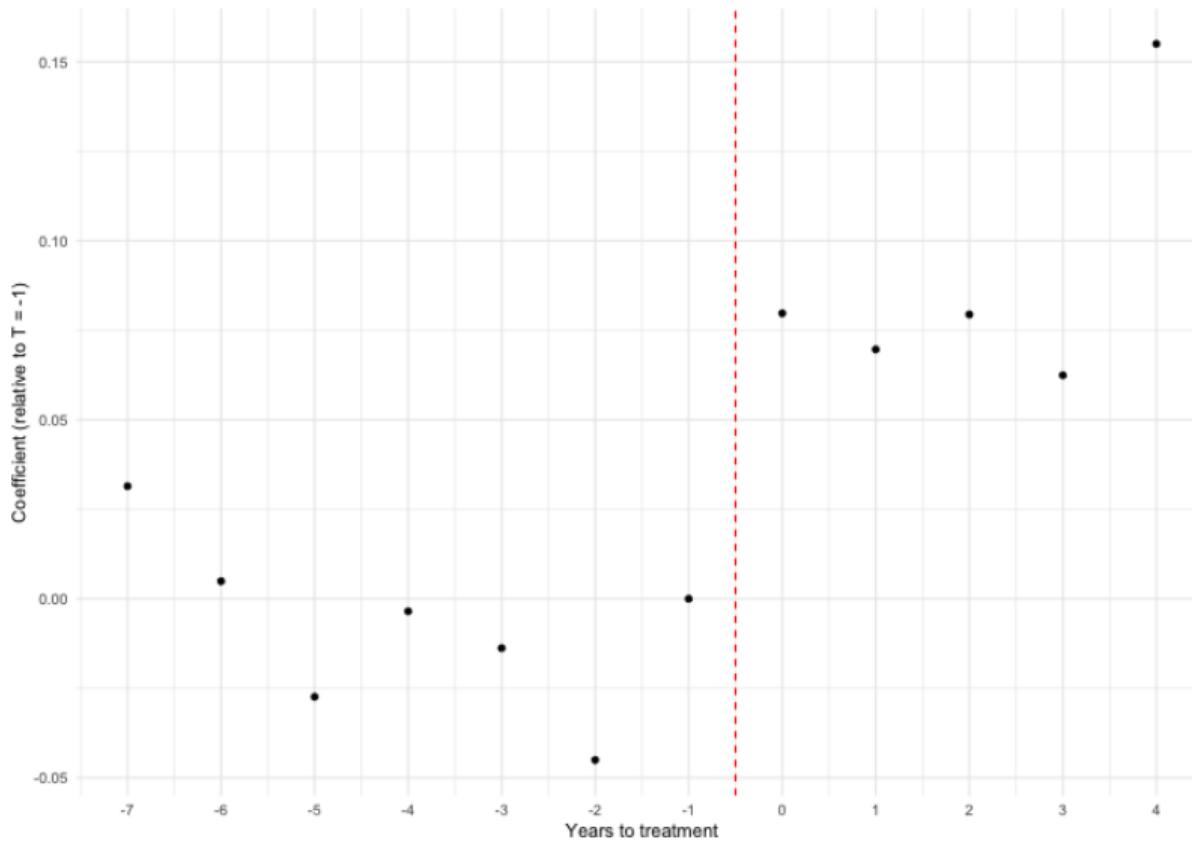
Plot these coefficients using geom_point

```
ggplot() +  
  geom_point(aes(x = years, y = coef)) +  
  geom_vline(xintercept = -0.5, linetype = "dashed", color = "red") +  
  labs(x = "Years to treatment", y = "Coefficient (relative to T = -1)") +  
  # change x axis to be every year  
  scale_x_continuous(breaks = years) +  
  theme_minimal()
```

Plot these coefficients using geom_point



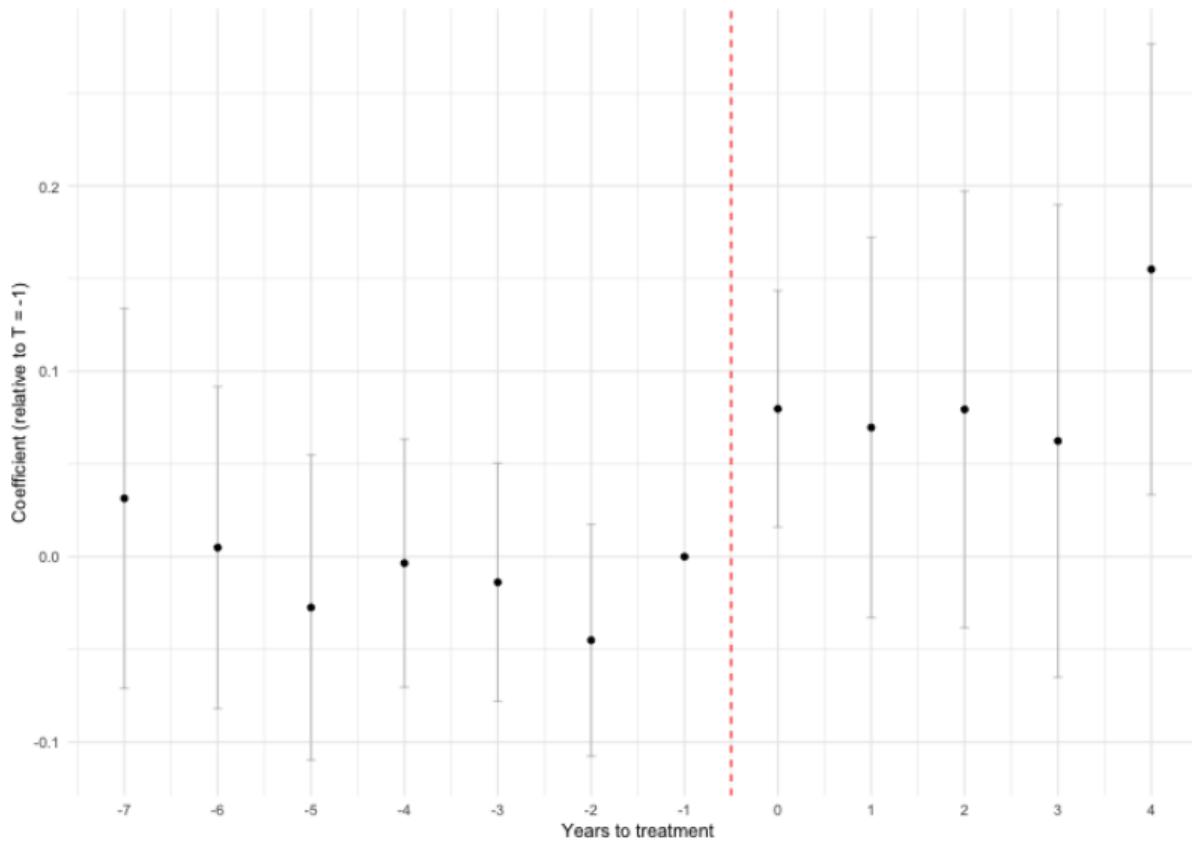
Let's remove the first three years because of small sample sizes



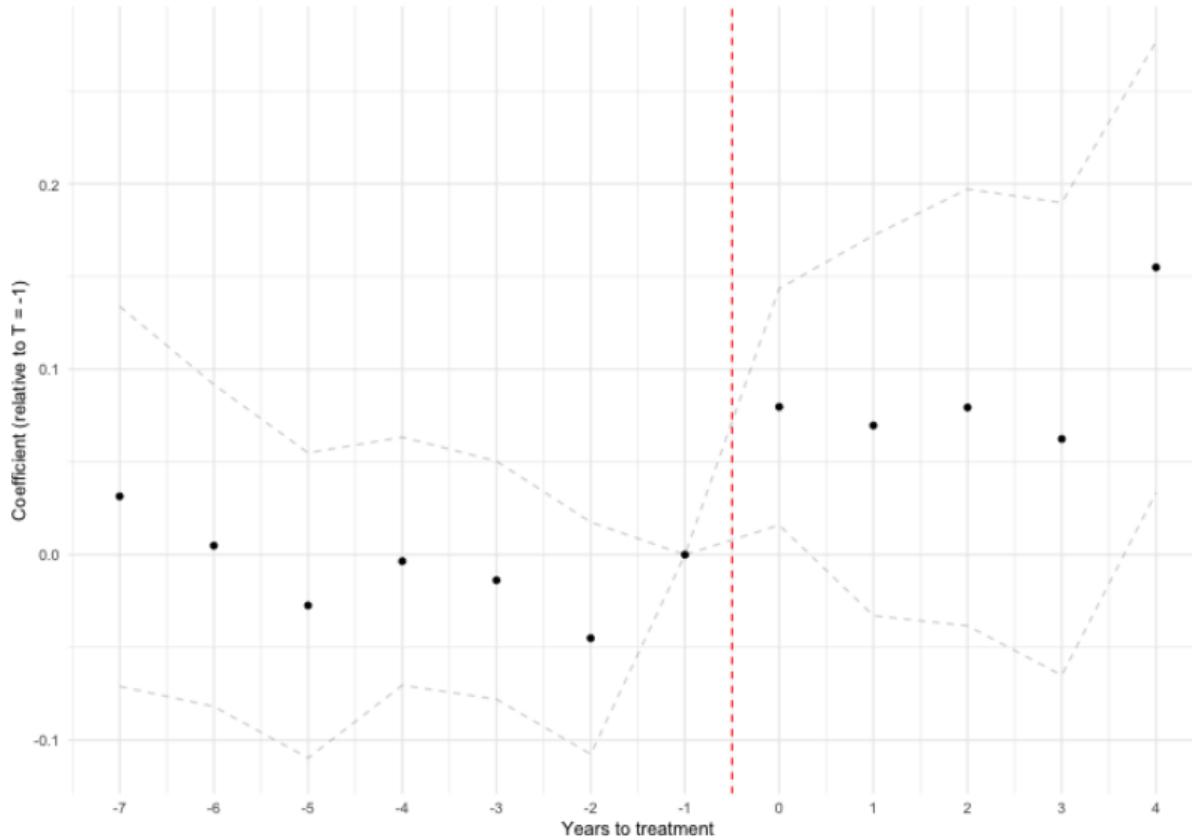
Add the confidence intervals using geom_errorbar

```
ggplot() +  
  geom_point(aes(x = years, y = coef)) +  
  geom_errorbar(aes(x = years, ymin = lower, ymax = upper), alpha = 0.2, width = 0.1) +  
  geom_vline(xintercept = -0.5, linetype = "dashed", color = "red") +  
  labs(x = "Years to treatment", y = "Coefficient (relative to T = -1)") +  
  # change x axis to be every year  
  scale_x_continuous(breaks = years) +  
  theme_minimal()
```

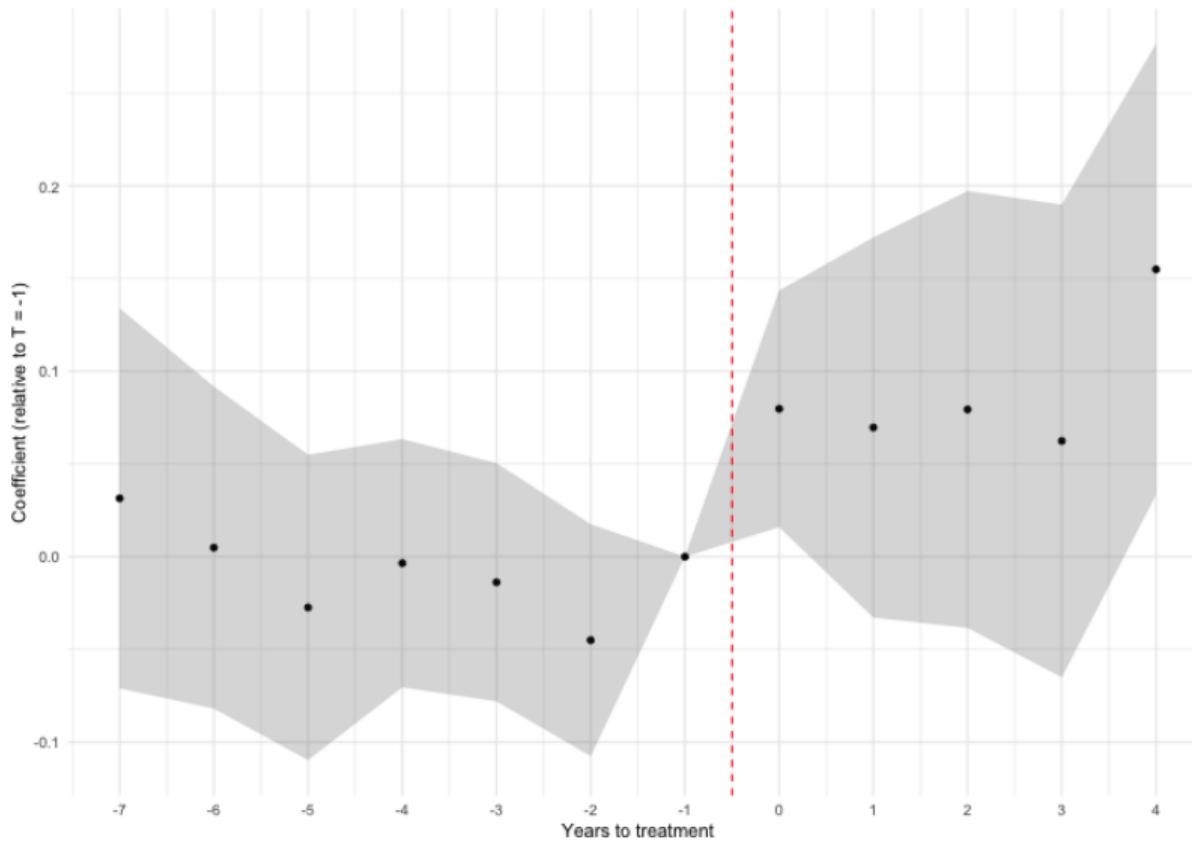
Add the confidence intervals using geom_errorbar



Add the confidence intervals using `geom_line`, if you prefer



Add the confidence intervals using `geom_ribbon`, if you prefer



Let's use the “Bacon decomposition” to look at weighting

- ▶ This won't allow us to calculate standard errors
- ▶ This of this as “diagnostics”
- ▶ We can see how different 2x2 cells have different weights, sometimes markedly so
 - ▶ We can also see that different cells have different treatment estimates

Let's try the “Bacon decomposition” - Note the treatment variable *must* be binary

```
library(bacondecomp)

# syntax:
# bacon(formula, data, id_var, time_var, quietly = F)
bacon <- bacon(homicide_c ~ cdl + log(population) + unemployrt, data = df, id_var = "state", time_var = "year", quietly = F)

##           type weight avg_est
## 1      Both Treated 0.10043 -0.00824
## 2 Treated vs Untreated 0.89957  0.07906
bacon$two_by_twos

##   treated untreated     weight     estimate          type
## 1    2007      99999 0.611262032  0.059301141 Treated vs Untreated
## 2    2007      2008 0.030235799 -0.001376830      Both Treated
## 3    2007      2010 0.017953382 -0.061066749      Both Treated
## 4    2007      2009 0.026493259  0.053198223      Both Treated
## 5    2006      2007 0.007511240  0.004557765      Both Treated
## 6    2006      99999 0.048373691  0.149043206 Treated vs Untreated
## 7    2006      2008 0.004157732  0.082718843      Both Treated
## 8    2006      2010 0.001570967 -0.048012581      Both Treated
## 9    2006      2009 0.002684252  0.086501842      Both Treated
## 10   2008      99999 0.160751757  0.092373673 Treated vs Untreated
## 11   2008      2010 0.004142796 -0.219504731      Both Treated
## 12   2008      2009 0.004503055 -0.151035583      Both Treated
## 13   2010      99999 0.016782728  0.076684232 Treated vs Untreated
## 14   2009      99999 0.062402696  0.184717537 Treated vs Untreated
## 15   2009      2010 0.001174611 -0.037899937      Both Treated
```

Get the overall average effect by multiplying weights by estimates

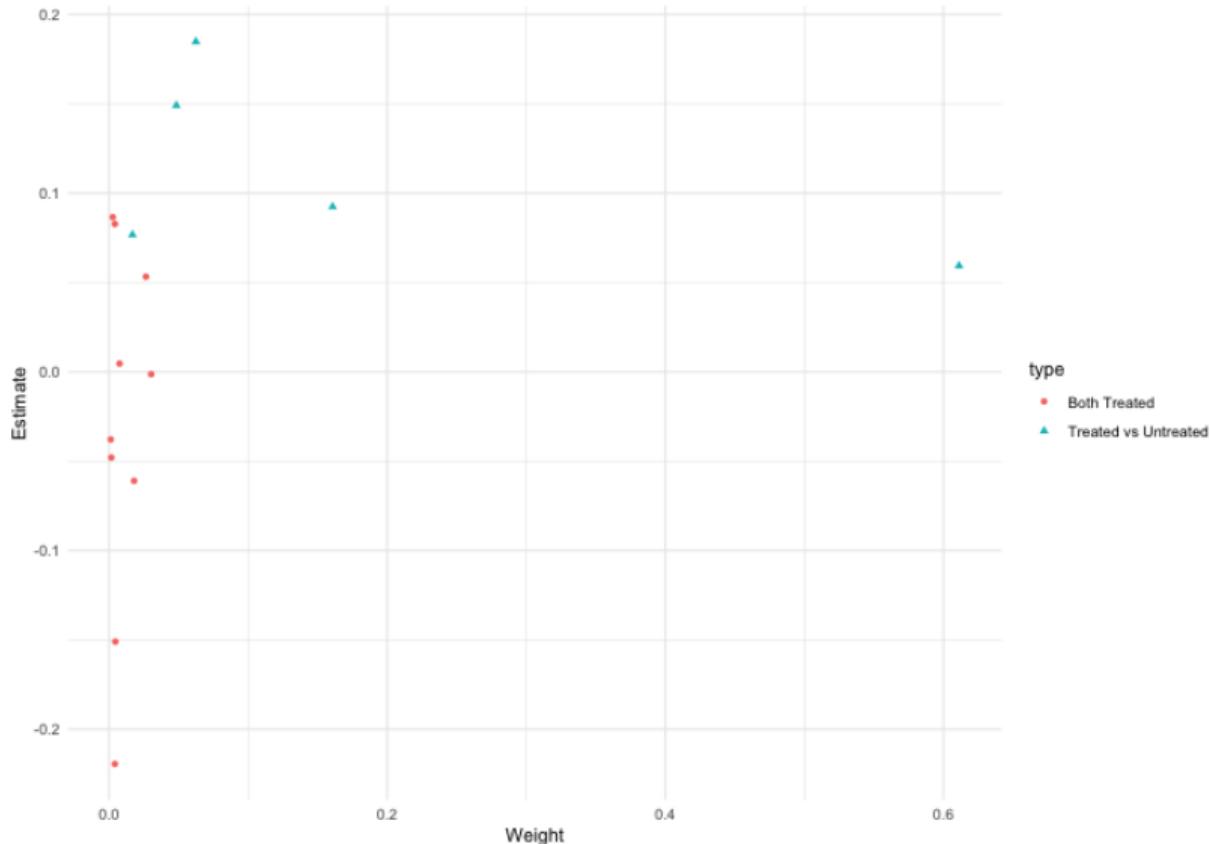
```
bacon <- bacon(homicide_c ~ cdl + log(population) + unemployrt, data = df, id_var = "state", time_var = "year", quietly = F)

##           type weight avg_est
## 1      Both Treated 0.10043 -0.00824
## 2 Treated vs Untreated 0.89957  0.07906
weighted.mean(bacon$two_by_twos$estimate, bacon$two_by_twos$weight)

## [1] 0.07029371
# compare to TWFE estimate
feols(homicide_c ~ cdl + log(population) + unemployrt | state + year, data = df,
      cluster = c("state")) # No weights since Bacon decomp doesn't allow them

## OLS estimation, Dep. Var.: homicide_c
## Observations: 550
## Fixed-effects: state: 50, year: 11
## Standard-errors: Clustered (state)
##           Estimate Std. Error   t value Pr(>|t|)
## cdl        0.073200  0.055921  1.308986 0.196645
## log(population) -1.018959  0.513389 -1.984769 0.052784 .
## unemployrt     -0.013406  0.013728 -0.976562 0.333583
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 0.175397   Adj. R2: 0.900242
## Within R2: 0.023689
```

We can plot the average effects for the different groups



Let's estimate effects using the did2s function from Kyle Butts

- ▶ bacondecomp is good for diagnostics, but we really want to estimate the ATT with standard errors

```
library(did2s)
```

- ▶ yname: the outcome variable
- ▶ first_stage: formula for first stage, can include fixed effects and covariates, but do not include treatment variable(s)!
- ▶ second_stage: This should be the treatment variable or in the case of event studies, treatment variables.
- ▶ treatment: This has to be the 0/1 treatment variable that marks when treatment turns on for a unit. If you suspect anticipation, see note above for accounting for this.
- ▶ cluster_var: Which variables to cluster on

Let's estimate effects using the did2s function from Kyle Butts

```
library(did2s)
# note that we can use fixest syntax, with FE and with i()!
# can also add weights
did2s <- did2s(data = df, yname = "homicide_c", first_stage = "log(population) + unemployrt | state + year",
                second_stage = "cdl", treatment = "cdl", cluster_var = "state", weights = "population")
# let's compare it to the vanilla TWFE
twfe <- feols(homicide_c ~ cdl + log(population) + unemployrt | state + year, data = df,
               cluster = c("state"), weights = df$population)
```

And the results

```
# The "correct" results
did2s

## OLS estimation, Dep. Var.: homicide_c
## Observations: 550
## Weights: weights_vector
## Standard-errors: Custom
##           Estimate Std. Error t value Pr(>|t|)
## cdl 0.094871   0.037989  2.4973 0.012806 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 262.4  Adj. R2: 0.085672

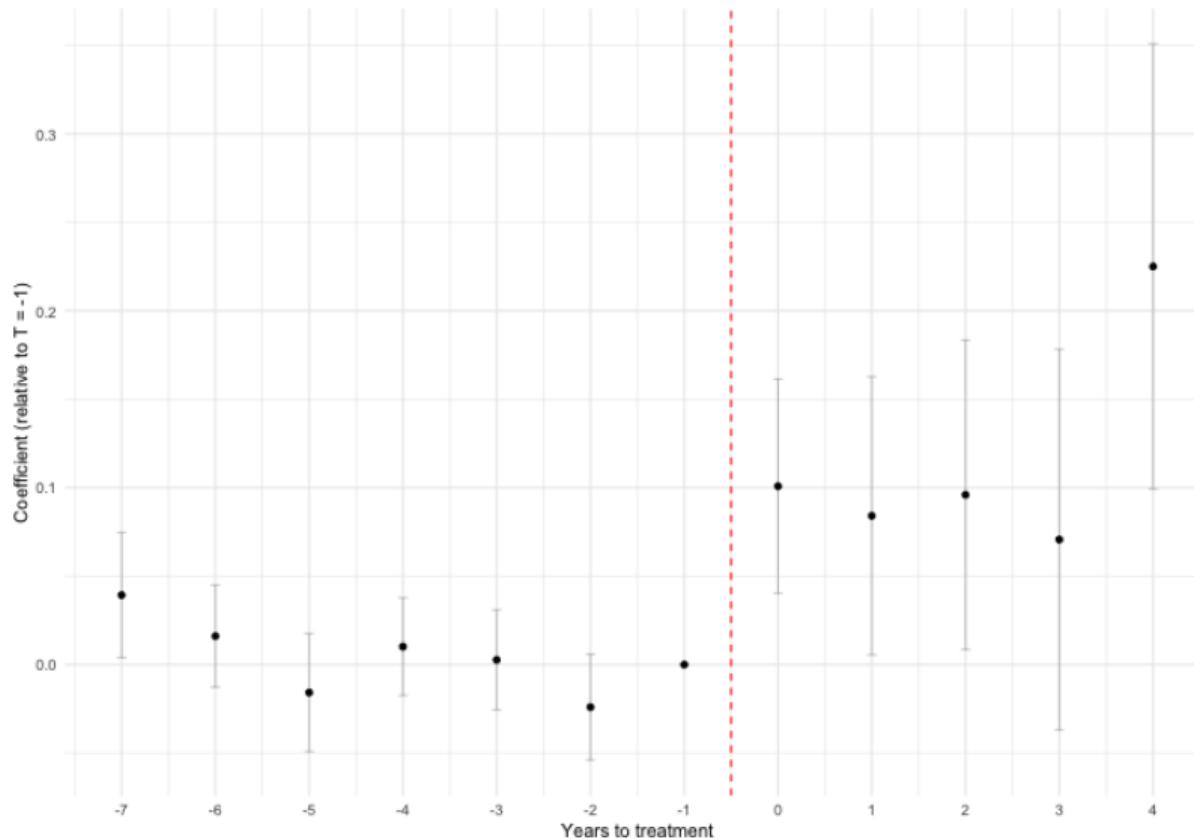
# and the TWFE results?
twfe

## OLS estimation, Dep. Var.: homicide_c
## Observations: 550
## Weights: df$population
## Fixed-effects: state: 50,  year: 11
## Standard-errors: Clustered (state)
##           Estimate Std. Error   t value Pr(>|t|)
## cdl          0.088939  0.032653  2.723783 0.0089189 **
## log(population) -0.767701  0.370173 -2.073895 0.0433658 *
## unemployrt     -0.009254  0.010207 -0.906587 0.3690621
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 252.9    Adj. R2: 0.927248
##           Within R2: 0.048831
```

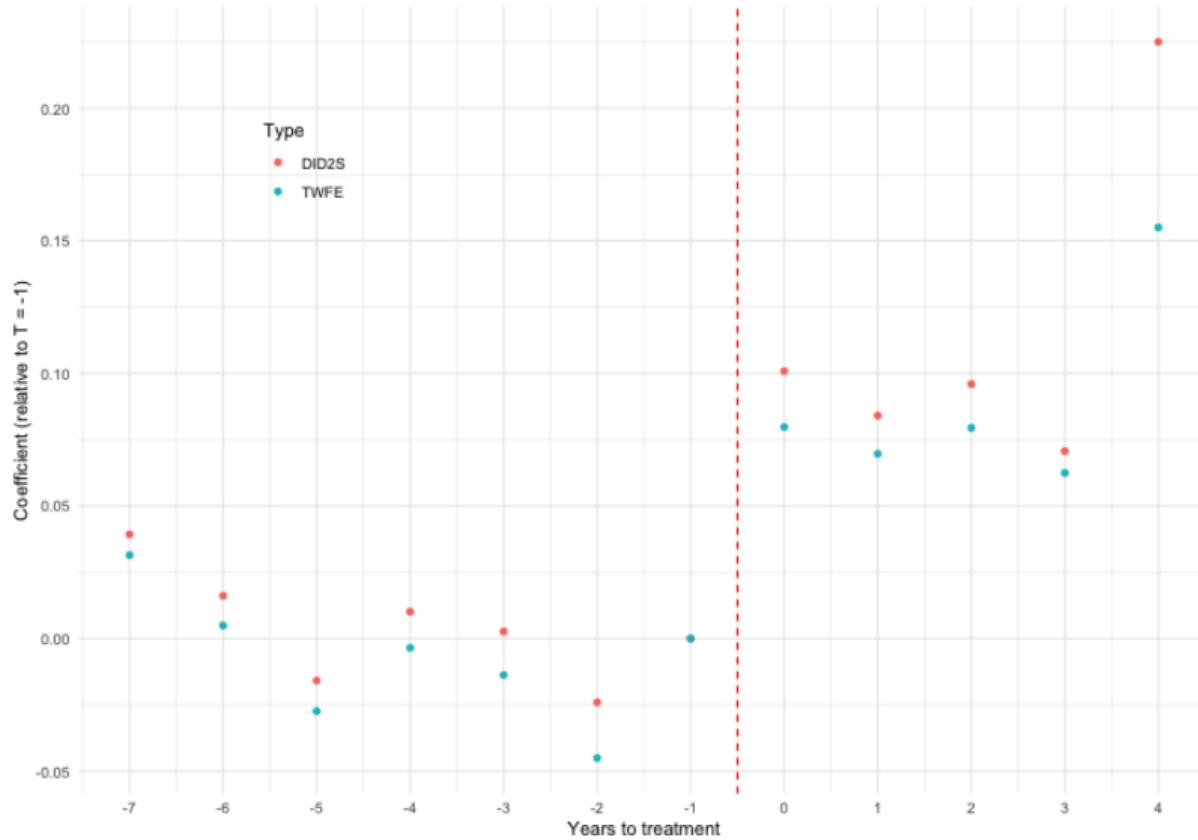
We can also estimate the event study this way!

```
# Estimate
did2s <- did2s(data = df, yname = "homicide_c", first_stage = "log(population) + unemployrt | state + year",
                 second_stage = "i(event_year, ref = -1)", treatment = "cdl", cluster_var = "state",
                 weights = "population")
coefficients <- c(did2s$coefficients[2:10], 0, did2s$coefficients[11:15])
# confidence intervals
lower <- c(confint(did2s)[2:10,1], 0, confint(did2s)[11:15,1])
upper <- c(confint(did2s)[2:10,2], 0, confint(did2s)[11:15,2])
# plot estimates
ggplot() +
  geom_point(aes(x = c(-10:4), y = coefficients)) +
  geom_errorbar(aes(x = c(-10:4), ymin = lower, ymax = upper), alpha = 0.2, width = 0.1) +
  geom_vline(xintercept = -0.5, linetype = "dashed", color = "red") +
  labs(x = "Years to treatment", y = "Coefficient (relative to T = -1)") +
# change x axis to be every year
  scale_x_continuous(breaks = years) +
  theme_minimal()
```

Some small changes to remove first three years



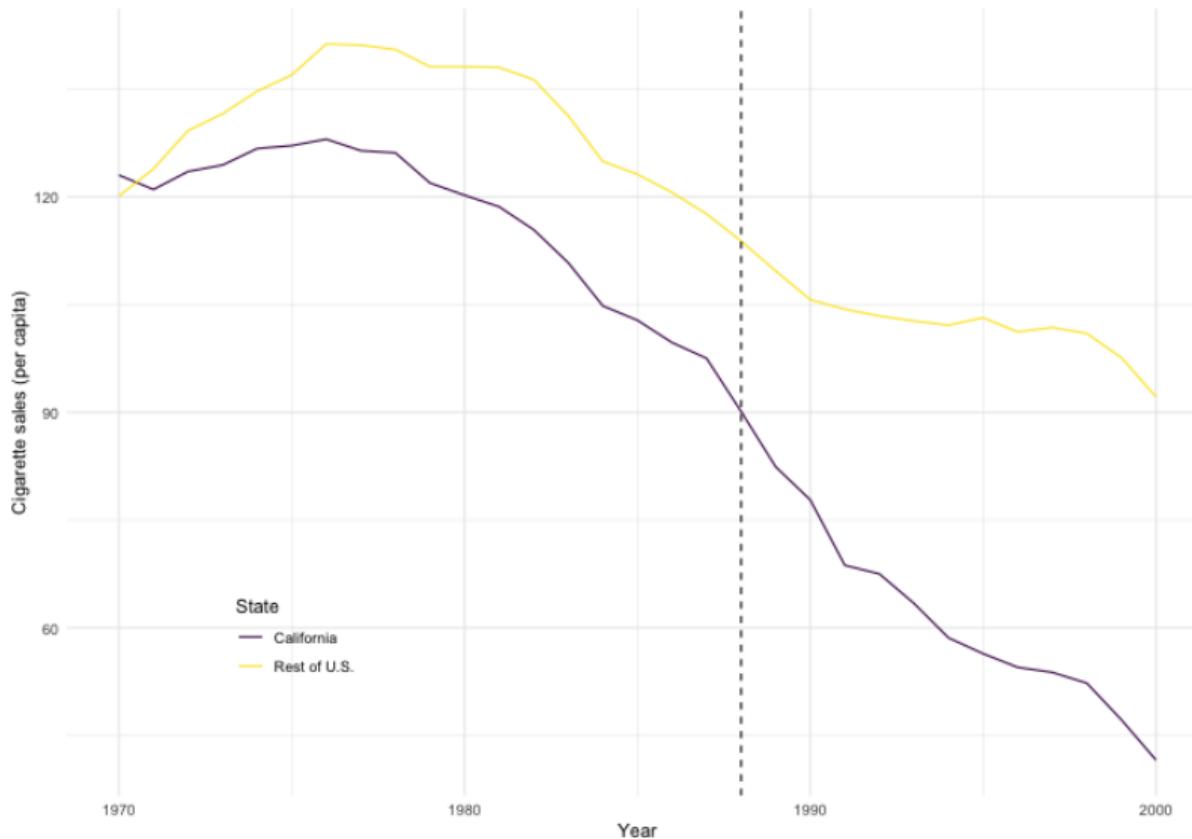
Compared to TWFE?



Wrapping up TWFE

- ▶ We've learned how to estimate two-way fixed effects models with `fixest`
 - ▶ We've also learned how they can be biased if treatment is staggered
- ▶ We saw how to use `did2s` to estimate the ATT
 - ▶ We also saw how to use `bacondecomp` to look at the weights
- ▶ A lingering question: TWFE with continuous treatment variables
 - ▶ Callaway et al. (2021) and Chaisemartin and D'Haultfœuille (2023) have some ideas
 - ▶ I haven't seen reliable packages for these, though
- ▶ IVs with TWFE?

Increasing cigarette taxes and discouraging smoking in California (in 1988)



Increasing cigarette taxes and discouraging smoking in California (in 1988)

- ▶ What if we wanted to figure out whether the law changed smoking in California?
- ▶ A key problem, similar to Card and Krueger:
 - ▶ There is really only one treated unit
 - ▶ Many people live in California, obviously, but they are all equally affected by the law
- ▶ Differences-in-differences is problematic here
 - ▶ A single treated cluster!
- ▶ So what are we to do?

Synthetic control

- ▶ This is where synthetic control comes in
 - ▶ Abadie and Gardeazabal (2003), *American Economic Review*
 - ▶ Abadie, Diamond, and Hainmueller (2010), *Journal of the American Statistical Association*
- ▶ The basic idea:
 - ▶ We want to create a “synthetic” California that is a weighted average of the other states
 - ▶ We want to weight the other states so that they look like California before the law
 - ▶ We can then compare the synthetic California to the real California after the law

Synthetic control

- ▶ Synthetic control doesn't just match on pre-treatment outcomes
 - ▶ It also matches on pre-treatment trends and covariates
 - ▶ This is what makes it different from pure matching
- ▶ Main requirement:
 - ▶ Many pre-treatment periods (Abadie, Diamond, and Hainmueller, 2010)

A little formalization

- ▶ Let's say we are interested in outcome Y_{jt} for unit j at time t
 - ▶ "Treated" group is $j = 1$
- ▶ In the post period, we estimate the effect of the intervention as

$$Y_{1t} - \sum_{j=1}^J w_j^* Y_{jt}, \quad (24)$$

where w_j are time-invariant weights that we will estimate in the *pre period*.

A little formalization

- ▶ In an ideal world, we would estimate the weights such that

$$\sum_{j=2}^J w_j^* Y_{j1} = Y_{11}, \sum_{j=2}^J w_j^* Y_{j2} = Y_{12}, \dots, \sum_{j=2}^J w_j^* Y_{jT_0} = Y_{1T_0}, \quad (25)$$

where T_0 is the number of pre-treatment periods.

- ▶ In practice, however, this will never hold exactly.
 - ▶ Instead, we will have to choose weights such that the differences are as small as possible.
- ▶ Our job is to estimate W , the vector of weights for each of the candidate control units.

A little formalization

- ▶ Consider a set of variables (which can include pre-treatment outcomes) X , where X_1 is the treated unit and X_0 are the control units.
- ▶ We will minimize

$$\|X_1 - X_0 W\| = \sqrt{(X_1 - X_0 W)' V (X_1 - X_0 W)}, \quad (26)$$

where V is a diagonal matrix of weights *for different variables*.

- ▶ Essentially there are two types of weights:
 - ▶ Weights for different variables
 - ▶ Weights for different units

A little formalization

- ▶ Estimating V is important
- ▶ The most common approach is to minimize mean squared prediction error (Cunnigham, 2022):

$$\sum_{t=1}^{T_0} \left(Y_{1t} - \sum_{j=1}^J w_j^* Y_{jt} \right)^2, \quad (27)$$

where, again, T_0 is the number of pre-treatment periods.

- ▶ Thankfully, the canned R packages do all of this for us!
 - ▶ It's nonetheless good to have an understanding of what they're doing

Synthetic control in R with the tidysynth package

```
library(tidysynth)

# We are going to use the "smoking" data from Abadie, Diamond, and Hainmueller (2010)
summary(df)

##      state      year     cigsale     lnincome       beer
## Min.   : 1   Min.   :1970   Min.   :40.7   Min.   : 9.397   Min.   : 2.50
## 1st Qu.:10  1st Qu.:1977  1st Qu.:100.9  1st Qu.: 9.739  1st Qu.:20.90
## Median :20  Median :1985  Median :116.3  Median : 9.861  Median :23.30
## Mean    :20  Mean   :1985  Mean   :118.9  Mean   : 9.862  Mean   :23.43
## 3rd Qu.:30  3rd Qu.:1993  3rd Qu.:130.5  3rd Qu.: 9.973  3rd Qu.:25.10
## Max.   :39   Max.   :2000   Max.   :296.2   Max.   :10.487  Max.   :40.40
## NA's    :195  NA's    :663
##      age15to24      retrprice     state_str      california
## Min.   :0.1294   Min.   : 27.3   Length:1209   Min.   :0.00000
## 1st Qu.:0.1658   1st Qu.: 50.0   Class :character 1st Qu.:0.00000
## Median :0.1781   Median : 95.5   Mode  :character Median :0.00000
## Mean   :0.1755   Mean   :108.3   NA's   :390      Mean   :0.02564
## 3rd Qu.:0.1867   3rd Qu.:158.4   NA's   :390      3rd Qu.:0.00000
## Max.   :0.2037   Max.   :351.2   NA's   :390      Max.   :1.00000
## NA's   :390
```

One issue: note how many missing values there are! We'll talk more on the next slide.

- ▶ A note: using this package is not straightforward. The website has code you can simply copy-paste:

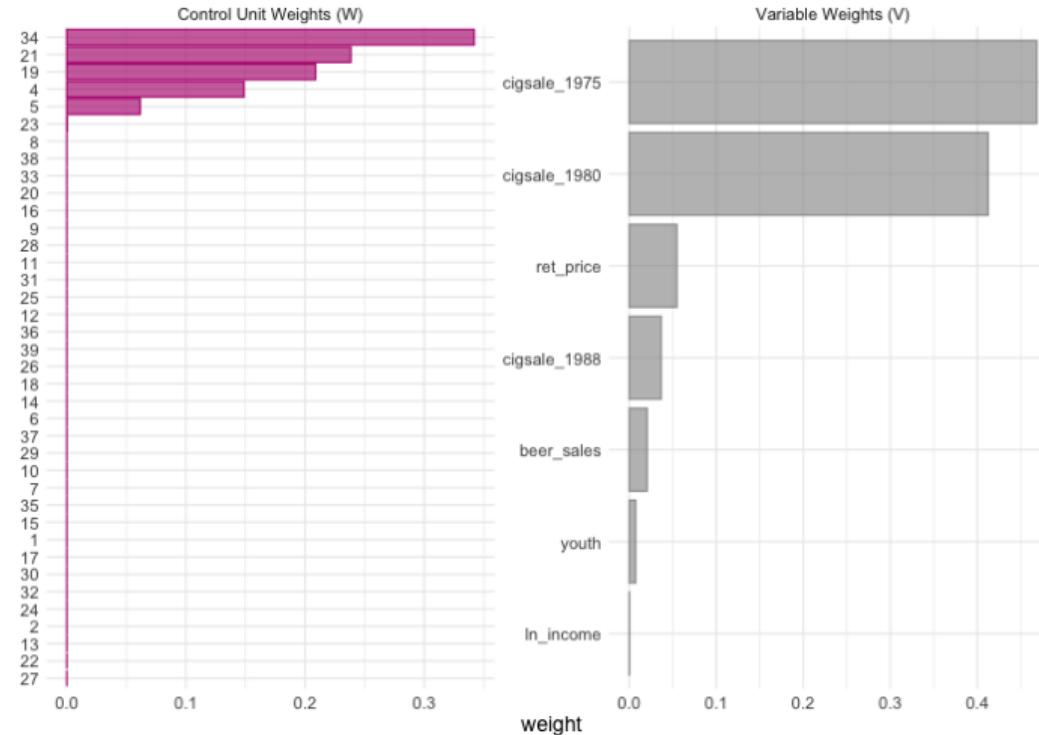
<https://cran.r-project.org/web/packages/tidysynth/readme/README.html>

We have to create a “synthetic control object”... kind of a pain, to be honest

```
smoking_out <- df %>%
  # initial the synthetic control object
  synthetic_control(outcome = cigsale, # outcome
                    unit = state, # unit index in the panel data
                    time = year, # time index in the panel data
                    i_unit = mean(df$state[df$california==1]), # unit where the intervention occurred
                    i_time = 1988, # time period when the intervention occurred
                    generate_placebos = T) %>% # generate placebo synthetic controls (for inference)
  # Generate the aggregate predictors used to fit the weights
  # average log income, retail price of cigarettes, and proportion of the
  # population between 15 and 24 years of age from 1980 - 1988
  generate_predictor(time_window = 1980:1988,
                      ln_income = mean(lnincome, na.rm = T),
                      ret_price = mean(retprice, na.rm = T),
                      youth = mean(age15to24, na.rm = T)) %>%
  # average beer consumption in the donor pool from 1984 - 1988
  generate_predictor(time_window = 1984:1988,
                      beer_sales = mean(beer, na.rm = T)) %>%
  # Lagged cigarette sales
  generate_predictor(time_window = 1975,
                     cigsale_1975 = cigsale) %>%
  generate_predictor(time_window = 1980,
                     cigsale_1980 = cigsale) %>%
  generate_predictor(time_window = 1988,
                     cigsale_1988 = cigsale) %>%
  # Generate the fitted weights for the synthetic control
  generate_weights(optimization_window = 1970:1988, # time to use in the optimization task
                  margin_ipop = .02, sigf_ipop = 7, bound_ipop = 6) %>% # optimizer options
  # Generate the synthetic control
  generate_control()
```

Let's first look at the weights

```
# If you get the first step above, the rest is easy  
smoking_out %>% plot_weights()
```



```
# but what are the states!?
```

Let's first look at the weights

```
# Grab the weights
weights <- smoking_out %>% grab_unit_weights()
# merge in state names
weights <- weights %>%
  mutate(unit = as.numeric(unit)) %>%
  left_join(df %>% select(unit = state, state_str) %>% distinct(),
            by = "unit")
# arrange by weight, in descending order
weights %>% arrange(-weight)
# note they sum to one
print(paste0("Sum of weights: ", sum(weights$weight)))
```

```
## # A tibble: 38 x 3
##       unit     weight state_str
##   <dbl>     <dbl> <chr>
## 1     34 0.342    Utah
## 2     21 0.238    Nevada
## 3     19 0.209    Montana
## 4      4 0.149    Colorado
## 5      5 0.0617   Connecticut
## 6     23 0.000434 New Mexico
## 7      8 0.00000341 Idaho
## 8     38 0.00000254 Wisconsin
## 9     33 0.00000163 Texas
## 10    20 0.00000153 Nebraska
## # i 28 more rows
## [1] "Sum of weights: 0.999999999995394"
```

Another common test: balance

```
balance <- smoking_out %>% grab_balance_table()
colnames(balance) <- c("Variable", "California", "Synthetic California", "Placebo")
kable(balance,
      align = "lccc", booktabs = TRUE, linesep = "", row.names = FALSE) %>%
  column_spec(1,width = "2cm") %>%
  column_spec(c(2:4),width = "2cm") %>%
  kable_styling()
```

Variable	California	Synthetic California	Placebo
ln_income	10.0765586	9.8534864	9.8291968
ret_price	89.4222234	89.3908482	87.2660819
youth	0.1735324	0.1736054	0.1725101
beer_sales	24.2800003	24.2224860	23.6552632
cigsale_1975	127.0999985	126.9914177	136.9315790
cigsale_1980	120.1999969	120.2204841	138.0894737
cigsale_1988	90.0999985	91.3888620	113.8236837

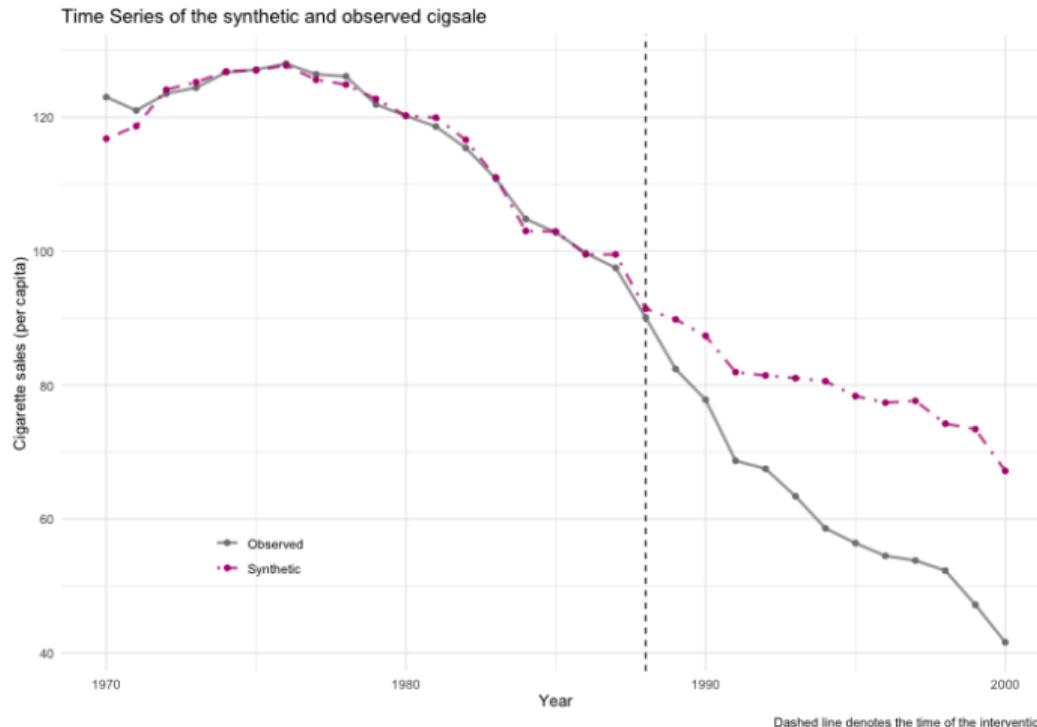
Another common test: balance

```
balance <- smoking_out %>% grab_balance_table()
colnames(balance) <- c("Variable", "California", "Synthetic California", "Placebo")
# let's round to make it prettier
balance[,2:4] <- round(balance[,2:4], 3)
kable(balance,
      align = "lccc", booktabs = TRUE, linesep = "", row.names = FALSE) %>%
  column_spec(1,width = "2cm") %>%
  column_spec(c(2:4),width = "2cm") %>%
  kable_styling()
```

Variable	California	Synthetic California	Placebo
ln_income	10.077	9.853	9.829
ret_price	89.422	89.391	87.266
youth	0.174	0.174	0.173
beer_sales	24.280	24.222	23.655
cigsale_1975	127.100	126.991	136.932
cigsale_1980	120.200	120.220	138.089
cigsale_1988	90.100	91.389	113.824

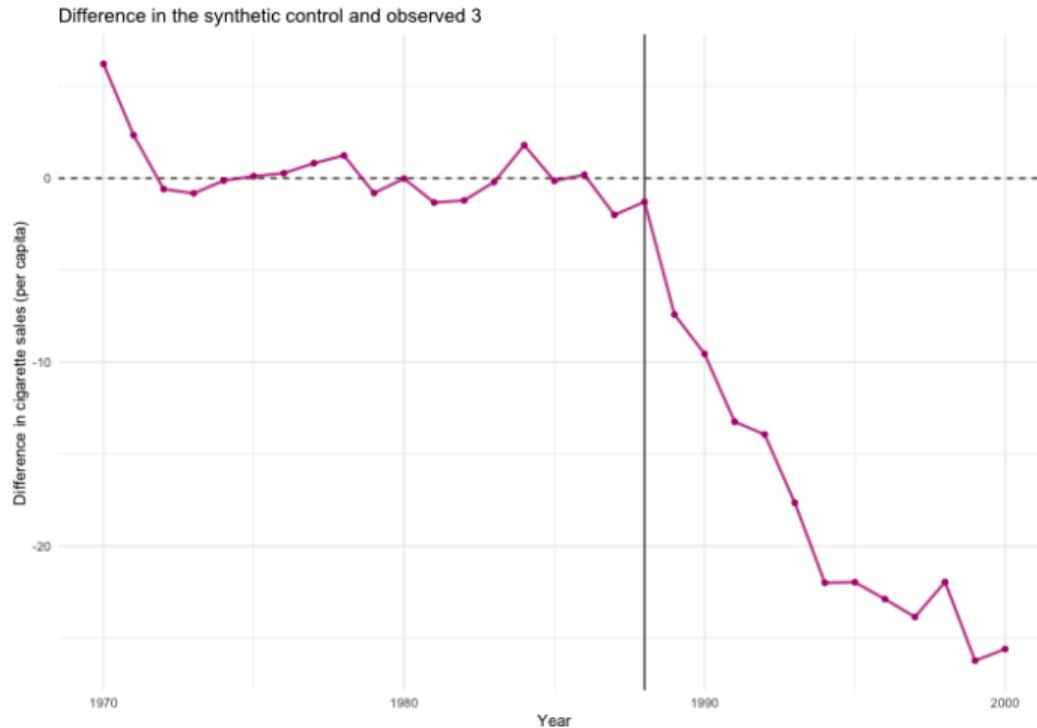
The results

```
gg1 <- smoking_out %>% plot_trends() # it's a ggplot object! So let's prettify it  
gg1 +  
  labs(x = "Year", y = "Cigarette sales (per capita)") +  
  theme(legend.position = c(0.2, 0.2))
```



Just the difference

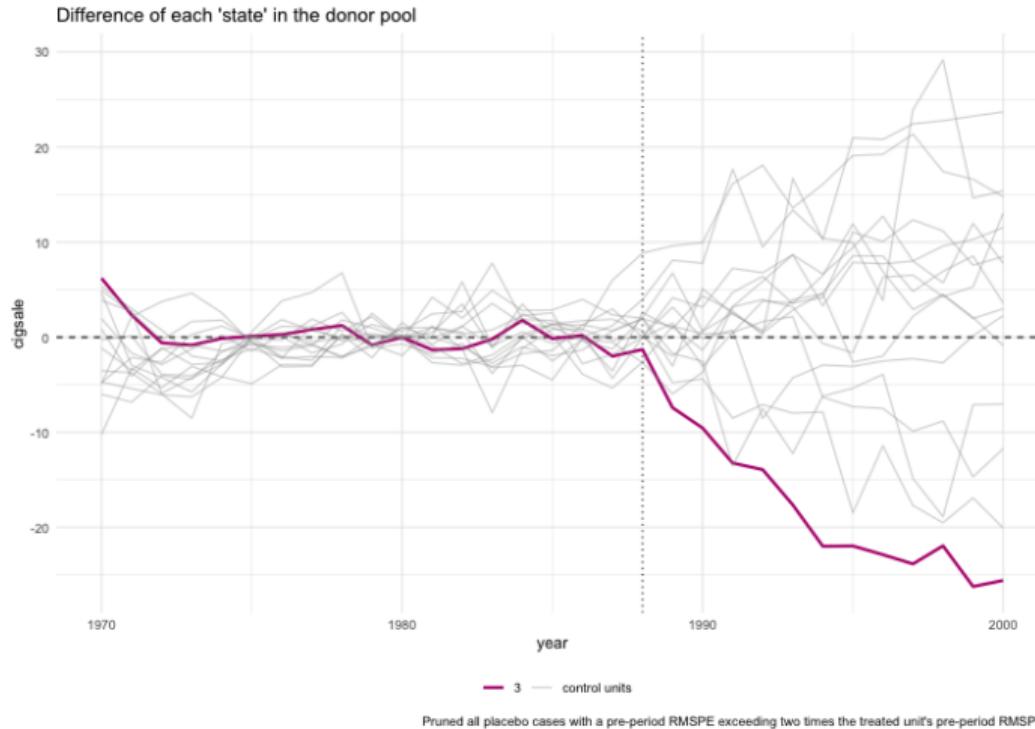
```
gg1 <- smoking_out %>% plot_differences()  
gg1 +  
  labs(x = "Year", y = "Difference in cigarette sales (per capita)")
```



- ▶ Okay... so is the result “big”?
 - ▶ I know what you’re really thinking: is the result *statistically significant*?
- ▶ There are no standard errors, as such, in synthetic control
- ▶ Instead, we use placebo tests
 - ▶ We basically do the exact same thing *for every state in our data*
 - ▶ If the effect is real, it should be much larger than any “effect” in other states, right?
 - ▶ `tidysynth` makes this easy

Placebo tests

```
smoking_out %>% plot_placebos()
```



Placebo tests

- ▶ It turns out what Abadie et al. (2010) do is a little more complicated
 - ▶ They look at mean squared prediction error (MSPE) for each state *before* and *after* the intervention
 - ▶ They then calculate the ratio: $\frac{MSPE_{after}}{MSPE_{before}}$
- ▶ We can rank states based on how much they *change* after the intervention!

Placebo tests

```
sig <- smoking_out %>% grab_significance()
# merge in state names
sig <- sig %>%
  mutate(unit = as.numeric(unit_name)) %>%
  left_join(df %>% select(unit = state, state_str) %>% distinct(), by = "unit")
# let's replace "unit_name" with "state_str"
sig$unit_name <- sig$state_str
sig <- sig %>% select(-c("state_str", "z_score", "unit"))
# round
sig[,c(3,4,5,7)] <- round(sig[,c(3,4,5,7)], 3)
kable(sig[1:10,], # just first 10
      align = "lcccccc", booktabs = TRUE, linesep = "", row.names = FALSE) %>%
  column_spec(c(1,7),width = "2cm") %>%
  column_spec(c(2:6),width = "1cm") %>%
  kable_styling()
```

unit_name	type	pre_mspe	post_mspe	mspe_ratio	rank	fishers_exact_pvalue
California	Treated	3.166	392.198	123.870	1	0.026
Georgia	Donor	3.786	178.712	47.208	2	0.051
Indiana	Donor	25.171	769.656	30.577	3	0.077
West Virginia	Donor	9.523	284.105	29.832	4	0.103
Wisconsin	Donor	11.134	267.763	24.050	5	0.128
Missouri	Donor	3.028	67.774	22.379	6	0.154
Texas	Donor	14.365	277.135	19.292	7	0.179
South Carolina	Donor	12.551	233.794	18.627	8	0.205
Virginia	Donor	9.809	96.397	9.828	9	0.231
Nebraska	Donor	6.297	52.891	8.400	10	0.256